

Appenders

Appenders are responsible for delivering LogEvents to their destination. Every Appender must implement the Appender (../log4j-core/apidocs/org/apache/logging/log4j/core/Appender.html) interface. Most Appenders will extend AbstractAppender (../log4j-core/apidocs/org/apache/logging/log4j/core/appender/AbstractAppender.html) which adds Lifecycle (../log4j-core/apidocs/org/apache/logging/log4j/core/Lifecycle.html) and Filterable (../log4j-core/apidocs/org/apache/logging/log4j/core/filter/Filterable.html) support. Lifecycle allows components to finish initialization after configuration has completed and to perform cleanup during shutdown. Filterable allows the component to have Filters attached to it which are evaluated during event processing.

Appenders usually are only responsible for writing the event data to the target destination. In most cases they delegate responsibility for formatting the event to a layout (layouts.html). Some appenders wrap other appenders so that they can modify the LogEvent, handle a failure in an Appender, route the event to a subordinate Appender based on advanced Filter criteria or provide similar functionality that does not directly format the event for viewing.

Appenders always have a name so that they can be referenced from Loggers.

In the tables below, the "Type" column corresponds to the Java type expected. For non-JDK classes, these should usually be in Log4j Core (../log4j-core/apidocs/index.html) unless otherwise noted.

AsyncAppender

The AsyncAppender accepts references to other Appenders and causes LogEvents to be written to them on a separate Thread. Note that exceptions while writing to those Appenders will be hidden from the application. The AsyncAppender should be configured after the appenders it references to allow it to shut down properly.

By default, AsyncAppender uses java.util.concurrent.ArrayBlockingQueue (<http://docs.oracle.com/javase/7/docs/api/java/util/concurrent/ArrayBlockingQueue.html>) which does not require any external libraries. Note that multi-threaded applications should exercise care when using this appender as such: the blocking queue is susceptible to lock contention and our tests showed (../performance.html#asyncLogging) performance may become worse when more threads are logging concurrently. Consider using lock-free Async Loggers (async.html) for optimal performance.

AsyncAppender Parameters		
Parameter Name	Type	Description
AppenderRef	String	The name of the Appenders to invoke asynchronously. Multiple AppenderRef elements can be configured.
blocking	boolean	If true, the appender will wait until there are free slots in the queue. If false, the event will be written to the error appender if the queue is full. The default is true.
shutdownTimeout	integer	How many milliseconds the Appender should wait to flush outstanding log events in the queue on shutdown. The default is zero which means to wait forever.
bufferSize	integer	Specifies the maximum number of events that can be queued. The default is 1024. Note that when using a disruptor-style BlockingQueue, this buffer size must be a power of 2. When the application is logging faster than the underlying appender can keep up with for a long enough time to fill up the queue, the behaviour is determined by the AsyncQueueFullPolicy (../log4j-core/apidocs/org/apache/logging/log4j/core/async/AsyncQueueFullPolicy.html).
errorRef	String	The name of the Appender to invoke if none of the appenders can be called, either due to errors in the appenders or because the queue is full. If not specified then errors will be ignored.
filter	Filter	A Filter to determine if the event should be handled by this Appender. More than one Filter may be used by using a CompositeFilter.
name	String	The name of the Appender.
ignoreExceptions	boolean	The default is true, causing exceptions encountered while appending events to be internally logged and then ignored. When set to false exceptions will be propagated to the caller, instead. You must set this to false when wrapping this Appender in a FailoverAppender.
includeLocation	boolean	Extracting location is an expensive operation (it can make logging 5 - 20 times slower). To improve performance, location is not included by default when adding a log event to the queue. You can change this by setting includeLocation="true".
BlockingQueueFactory	BlockingQueueFactory	This element overrides what type of BlockingQueue to use. See below documentation for more details.

There are also a few system properties that can be used to maintain application throughput even when the underlying appender cannot keep up with the logging rate and the queue is filling up. See the details for system properties log4j2.AsyncQueueFullPolicy and log4j2.DiscardThreshold (configuration.html#log4j2.AsyncQueueFullPolicy).

A typical AsyncAppender configuration might look like:

```
1. <?xml version="1.0" encoding="UTF-8"?>
2. <Configuration status="warn" name="MyApp" packages="">
3.   <Appenders>
4.     <File name="MyFile" fileName="logs/app.log">
5.       <PatternLayout>
6.         <Pattern>%d %p %c{1.} [%t] %m%n</Pattern>
7.       </PatternLayout>
8.     </File>
9.     <Async name="Async">
10.      <AppenderRef ref="MyFile"/>
11.    </Async>
12.  </Appenders>
13.  <Loggers>
14.    <Root level="error">
15.      <AppenderRef ref="Async"/>
16.    </Root>
17.  </Loggers>
18. </Configuration>
```

Starting in Log4j 2.7, a custom implementation of BlockingQueue or TransferQueue can be specified using a BlockingQueueFactory ([../log4j-core/apidocs/org/apache/logging/log4j/core/async/BlockingQueueFactory.html](#)) plugin. To override the default BlockingQueueFactory, specify the plugin inside an <Async/> element like so:

```
1. <Configuration name="LinkedTransferQueueExample">
2.   <Appenders>
3.     <List name="List"/>
4.     <Async name="Async" bufferSize="262144">
5.       <AppenderRef ref="List"/>
6.       <LinkedTransferQueue/>
7.     </Async>
8.   </Appenders>
9.   <Loggers>
10.    <Root>
11.      <AppenderRef ref="Async"/>
12.    </Root>
13.  </Loggers>
14. </Configuration>
```

Log4j ships with the following implementations:

BlockingQueueFactory Implementations

Plugin Name	Description
ArrayBlockingQueue	This is the default implementation that uses ArrayBlockingQueue (https://docs.oracle.com/javase/7/docs/api/java/util/concurrent/ArrayBlockingQueue.html).
DisruptorBlockingQueue	This uses the Conversant Disruptor (https://github.com/conversant/disruptor) implementation of BlockingQueue. This plugin takes a single optional attribute, spinPolicy, which corresponds to
JCToolsBlockingQueue	This uses JCTools (https://jctools.github.io/JCTools/), specifically the <u>MPSC (multiple producer single consumer)</u> bounded lock-free queue.
LinkedTransferQueue	This uses the new in Java 7 implementation LinkedTransferQueue (https://docs.oracle.com/javase/7/docs/api/java/util/concurrent/LinkedTransferQueue.html). Note that this queue does not use the bufferSize configuration attribute from AsyncAppender as LinkedTransferQueue does not support a maximum capacity.

CassandraAppender

The CassandraAppender writes its output to an Apache Cassandra (<https://cassandra.apache.org/>) database. A keyspace and table must be configured ahead of time, and the columns of that table are mapped in a configuration file. Each column can specify either a StringLayout (layouts.html) (e.g., a PatternLayout (layouts.html#PatternLayout)) along with an optional conversion type, or only a conversion type for org.apache.logging.log4j.spi.ThreadContextMap or org.apache.logging.log4j.spi.ThreadContextStack to store the MDC or NDC (thread-context.html) in a map or list column respectively. A conversion type compatible with java.util.Date will use the log event timestamp converted to that type (e.g., use java.util.Date to fill a timestamp column type in Cassandra).

CassandraAppender Parameters

Parameter Name	Type	Description
batched	boolean	Whether or not to use batch statements false.
batchType	BatchStatement.Type (http://docs.datastax.com/en/drivers/java/3.0/com/datastax/driver/core/BatchStatement.Type.html)	The batch type to use when using batched
bufferSize	int	The number of log messages to buffer in
clusterName	String	The name of the Cassandra cluster to c

columns	ColumnMapping[]	A list of column mapping configurations can have a conversion type specified by the column name. If the configured type is a String, then the log timestamp attribute is given, then its value will be used. Otherwise, the layout or pattern specific that column.
contactPoints	SocketAddress[]	A list of hosts and ports of Cassandra replication addresses. By default, if a port is not specified, the default Cassandra port of 9042 will be used. By default, this is empty.
filter	Filter	A Filter to determine if the event should be used by using a CompositeFilter.
ignoreExceptions	boolean	The default is true, causing exceptions logged and then ignored. When set to false, you must set this to false when wrapping the logger.
keyspace	String	The name of the keyspace containing the table.
name	String	The name of the Appender.
password	String	The password to use (along with the username) to connect to Cassandra.
table	String	The name of the table to write log messages to.
useClockForTimestampGenerator	boolean	Whether or not to use the configured or default TimestampGenerator (http://docs.datastax.com/en/drivers/java/2.1.0/index.html). By default, this is false.
username	String	The username to use to connect to Cassandra.
useTls	boolean	Whether or not to use TLS/SSL to connect to Cassandra.

Here is an example CassandraAppender configuration:

```
1. <Configuration name="CassandraAppenderTest">
2.   <Appenders>
3.     <Cassandra name="Cassandra" clusterName="Test Cluster" keyspace="test" table="logs" bufferSize="10" batched="true">
4.       <SocketAddress host="localhost" port="9042"/>
5.       <ColumnMapping name="id" pattern="%uuid{TIME}" type="java.util.UUID"/>
6.       <ColumnMapping name="timeid" literal="now()"/>
7.       <ColumnMapping name="message" pattern="%message"/>
8.       <ColumnMapping name="level" pattern="%level"/>
9.       <ColumnMapping name="marker" pattern="%marker"/>
10.      <ColumnMapping name="logger" pattern="%logger"/>
11.      <ColumnMapping name="timestamp" type="java.util.Date"/>
12.      <ColumnMapping name="mdc" type="org.apache.logging.log4j.spi.ThreadContextMap"/>
13.      <ColumnMapping name="ndc" type="org.apache.logging.log4j.spi.ThreadContextStack"/>
14.    </Cassandra>
15.  </Appenders>
16.  <Loggers>
17.    <Logger name="org.apache.logging.log4j.cassandra" level="DEBUG">
18.      <AppenderRef ref="Cassandra"/>
19.    </Logger>
20.    <Root level="ERROR"/>
21.  </Loggers>
22. </Configuration>
```

This example configuration uses the following table schema:

```
1. CREATE TABLE logs (
2.   id timeuuid PRIMARY KEY,
3.   timeid timeuuid,
4.   message text,
5.   level text,
6.   marker text,
7.   logger text,
8.   timestamp timestamp,
9.   mdc map<text,text>,
10.  ndc list<text>
11. );
```

ConsoleAppender

As one might expect, the ConsoleAppender writes its output to either System.out or System.err with System.out being the default target. A Layout must be provided to format the LogEvent.

ConsoleAppender Parameters		
Parameter Name	Type	Description
filter	Filter	A Filter to determine if the event should be handled by this Appender. More than one Filter may be used by using a CompositeFilter.
layout	Layout	The Layout to use to format the LogEvent. If no layout is supplied the default pattern layout of "%m%n" will be used.
follow	boolean	Identifies whether the appender honors reassignments of System.out or System.err via System.setOut or System.setErr made after configuration. Note that the follow attribute cannot be used with Jansi on Windows. Cannot be used with direct.
direct	boolean	Write directly to java.io.FileDescriptor and bypass java.lang.System.out/.err. Can give up to 10x performance boost when the output is redirected to file or other process. Cannot be used with Jansi on Windows. Cannot be used with follow. Output will not respect java.lang.System.setOut()/setErr() and may get intertwined with other output to java.lang.System.out/.err in a multi-threaded application. <i>New since 2.6.2. Be aware that this is a new addition, and it has only been tested with Oracle JVM on Linux and Windows so far.</i>
name	String	The name of the Appender.
ignoreExceptions	boolean	The default is true, causing exceptions encountered while appending events to be internally logged and then ignored. When set to false exceptions will be propagated to the caller, instead. You must set this to false when wrapping this Appender in a FailoverAppender.
target	String	Either "SYSTEM_OUT" or "SYSTEM_ERR". The default is "SYSTEM_OUT".

A typical Console configuration might look like:

```
1. <?xml version="1.0" encoding="UTF-8"?>
2. <Configuration status="warn" name="MyApp" packages="">
3.   <Appenders>
4.     <Console name="STDOUT" target="SYSTEM_OUT">
5.       <PatternLayout pattern="%m%n" />
6.     </Console>
7.   </Appenders>
8.   <Loggers>
9.     <Root level="error">
10.      <AppenderRef ref="STDOUT" />
11.    </Root>
12.  </Loggers>
13. </Configuration>
```

FailoverAppender

The FailoverAppender wraps a set of appenders. If the primary Appender fails the secondary appenders will be tried in order until one succeeds or there are no more secondaries to try.

FailoverAppender Parameters		
Parameter Name	Type	Description
filter	Filter	A Filter to determine if the event should be handled by this Appender. More than one Filter may be used by using a CompositeFilter.
primary	String	The name of the primary Appender to use.
failovers	String[]	The names of the secondary Appenders to use.
name	String	The name of the Appender.
retryIntervalSeconds	integer	The number of seconds that should pass before retrying the primary Appender. The default is 60.
ignoreExceptions	boolean	The default is true, causing exceptions encountered while appending events to be internally logged and then ignored. When set to false exceptions will be propagated to the caller, instead.
target	String	Either "SYSTEM_OUT" or "SYSTEM_ERR". The default is "SYSTEM_ERR".

A Failover configuration might look like:

```
1. <?xml version="1.0" encoding="UTF-8"?>
2. <Configuration status="warn" name="MyApp" packages="">
3.   <Appenders>
4.     <RollingFile name="RollingFile" fileName="logs/app.log" filePattern="logs/app-%d{MM-dd-yyyy}.log.gz"
5.       ignoreExceptions="false">
6.       <PatternLayout>
7.         <Pattern>%d %p %c{1.} [%t] %m%n</Pattern>
8.       </PatternLayout>
9.       <TimeBasedTriggeringPolicy />
10.    </RollingFile>
11.    <Console name="STDOUT" target="SYSTEM_OUT" ignoreExceptions="false">
12.      <PatternLayout pattern="%m%n"/>
13.    </Console>
14.    <Failover name="Failover" primary="RollingFile">
15.      <Failovers>
16.        <AppenderRef ref="Console"/>
17.      </Failovers>
18.    </Failover>
19.  </Appenders>
20.  <Loggers>
21.    <Root level="error">
22.      <AppenderRef ref="Failover"/>
23.    </Root>
24.  </Loggers>
25. </Configuration>
```

FileAppender

The FileAppender is an OutputStreamAppender that writes to the File named in the fileName parameter. The FileAppender uses a FileManager (which extends OutputStreamManager) to actually perform the file I/O. While FileAppenders from different Configurations cannot be shared, the FileManagers can be if the Manager is accessible. For example, two web applications in a servlet container can have their own configuration and safely write to the same file if Log4j is in a ClassLoader that is common to both of them.

FileAppender Parameters

Parameter Name	Type	Description
append	boolean	When true - the default, records will be appended to the end of the file. When set to false, the file will be cleared before new records are written.
bufferedIO	boolean	When true - the default, records will be written to a buffer and the data will be written to disk when the buffer is full or, if immediateFlush is set, when the record is written. File locking cannot be used with bufferedIO. Performance tests have shown that using buffered I/O significantly improves performance, even if immediateFlush is enabled.
bufferSize	int	When bufferedIO is true, this is the buffer size, the default is 8192 bytes.
createOnDemand	boolean	The appender creates the file on-demand. The appender only creates the file when a log event passes all filters and is routed to this appender. Defaults to false.
filter	Filter	A Filter to determine if the event should be handled by this Appender. More than one Filter may be used by using a CompositeFilter.
fileName	String	The name of the file to write to. If the file, or any of its parent directories, do not exist, they will be created.
immediateFlush	boolean	When set to true - the default, each write will be followed by a flush. This will guarantee the data is written to disk but could impact performance. Flushing after every write is only useful when using this appender with synchronous loggers. Asynchronous loggers and appenders will automatically flush at the end of a batch of events, even if immediateFlush is set to false. This also guarantees the data is written to disk but is more efficient.
layout	Layout	The Layout to use to format the LogEvent. If no layout is supplied the default pattern layout of "%m%n" will be used.
locking	boolean	When set to true, I/O operations will occur only while the file lock is held allowing FileAppenders in multiple JVMs and potentially multiple hosts to write to the same file simultaneously. This will significantly impact performance so should be used carefully. Furthermore, on many systems the file lock is "advisory" meaning that other applications can perform operations on the file without acquiring a lock. The default value is false.
name	String	The name of the Appender.
ignoreExceptions	boolean	The default is true, causing exceptions encountered while appending events to be internally logged and then ignored. When set to false exceptions will be propagated to the caller, instead. You must set this to false when wrapping this Appender in a FailoverAppender.
filePermissions	String	File attribute permissions in POSIX format to apply whenever the file is created. Underlying files system shall support POSIX (https://docs.oracle.com/javase/7/docs/api/java/nio/file/attribute/PosixFileAttributeView.html) file attribute view. Examples: rw----- or rw-rw-rw- etc...

fileOwner	String	<p>File owner to define whenever the file is created.</p> <p>Changing file's owner may be restricted for security reason and Operation not permitted IOException thrown. Only processes with an effective user ID equal to the user ID of the file or with appropriate privileges may change the ownership of a file if <code>_POSIX_CHOWN_RESTRICTED</code> (http://www.gnu.org/software/libc/manual/html_node/Options-for-Files.html) is in effect for path.</p> <p>Underlying files system shall support file owner (https://docs.oracle.com/javase/7/docs/api/java/nio/file/attribute/FileOwnerAttributeView.html) attribute view.</p>
fileGroup	String	<p>File group to define whenever the file is created.</p> <p>Underlying files system shall support POSIX (https://docs.oracle.com/javase/7/docs/api/java/nio/file/attribute/PosixFileAttributeView.html) file attribute view.</p>

Here is a sample File configuration:

```
1. <?xml version="1.0" encoding="UTF-8"?>
2. <Configuration status="warn" name="MyApp" packages="">
3.   <Appenders>
4.     <File name="MyFile" fileName="logs/app.log">
5.       <PatternLayout>
6.         <Pattern>%d %p %c{1.} [%t] %m%n</Pattern>
7.       </PatternLayout>
8.     </File>
9.   </Appenders>
10.  <Loggers>
11.    <Root level="error">
12.      <AppenderRef ref="MyFile"/>
13.    </Root>
14.  </Loggers>
15. </Configuration>
```

FlumeAppender

This is an optional component supplied in a separate jar.

Apache Flume (<http://flume.apache.org/index.html>) is a distributed, reliable, and available system for efficiently collecting, aggregating, and moving large amounts of log data from many different sources to a centralized data store. The FlumeAppender takes LogEvents and sends them to a Flume agent as serialized Avro events for consumption.

The Flume Appender supports three modes of operation.

1. It can act as a remote Flume client which sends Flume events via Avro to a Flume Agent configured with an Avro Source.
2. It can act as an embedded Flume Agent where Flume events pass directly into Flume for processing.
3. It can persist events to a local BerkeleyDB data store and then asynchronously send the events to Flume, similar to the embedded Flume Agent but without most of the Flume dependencies.

Usage as an embedded agent will cause the messages to be directly passed to the Flume Channel and then control will be immediately returned to the application. All interaction with remote agents will occur asynchronously. Setting the "type" attribute to "Embedded" will force the use of the embedded agent. In addition, configuring agent properties in the appender configuration will also cause the embedded agent to be used.

FlumeAppender Parameters		
Parameter Name	Type	Description
agents	Agent[]	An array of Agents to which the logging events should be sent. If more than one agent is specified the first Agent will be the primary and subsequent Agents will be used in the order specified as secondaries should the primary Agent fail. Each Agent definition supplies the Agents host and port. The specification of agents and properties are mutually exclusive. If both are configured an error will result.
agentRetries	integer	The number of times the agent should be retried before failing to a secondary. This parameter is ignored when type="persistent" is specified (agents are tried once before failing to the next).
batchSize	integer	Specifies the number of events that should be sent as a batch. The default is 1. <i>This parameter only applies to the Flume Appender.</i>
compress	boolean	When set to true the message body will be compressed using gzip
connectTimeoutMillis	integer	The number of milliseconds Flume will wait before timing out the connection.
dataDir	String	Directory where the Flume write ahead log should be written. Valid only when embedded is set to true and Agent elements are used instead of Property elements.
filter	Filter	A Filter to determine if the event should be handled by this Appender. More than one Filter may be used by using a CompositeFilter.
eventPrefix	String	The character string to prepend to each event attribute in order to distinguish it from MDC attributes. The default is an empty string.
flumeEventFactory	FlumeEventFactory	Factory that generates the Flume events from Log4j events. The default factory is the FlumeAvroAppender itself.
layout	Layout	The Layout to use to format the LogEvent. If no layout is specified RFC5424Layout will be used.
lockTimeoutRetries	integer	The number of times to retry if a LockConflictException occurs while writing to Berkeley DB. The default is 5.
maxDelayMillis	integer	The maximum number of milliseconds to wait for batchSize events before publishing the batch.

mdcExcludes	String	A comma separated list of mdc keys that should be excluded from the FlumeEvent. This is mutually exclusive with the mdcIncludes attribute.
mdcIncludes	String	A comma separated list of mdc keys that should be included in the FlumeEvent. Any keys in the MDC not found in the list will be excluded. This option is mutually exclusive with the mdcExcludes attribute.
mdcRequired	String	A comma separated list of mdc keys that must be present in the MDC. If a key is not present a LoggingException will be thrown.
mdcPrefix	String	A string that should be prepended to each MDC key in order to distinguish it from event attributes. The default string is "mdc:".
name	String	The name of the Appender.
properties	Property[]	<p>One or more Property elements that are used to configure the Flume Agent. The properties must be configured without the agent name (the appender name is used for this) and no sources can be configured. Interceptors can be specified for the source using "sources.log4j-source.interceptors". All other Flume configuration properties are allowed. Specifying both Agent and Property elements will result in an error.</p> <p>When used to configure in Persistent mode the valid properties are:</p> <ol style="list-style-type: none">1. "keyProvider" to specify the name of the plugin to provide the secret key for encryption.
requestTimeoutMillis	integer	The number of milliseconds Flume will wait before timing out the request.
ignoreExceptions	boolean	The default is true, causing exceptions encountered while appending events to be internally logged and then ignored. When set to false exceptions will be propagated to the caller, instead. You must set this to false when wrapping this Appender in a FailoverAppender.
type	enumeration	One of "Avro", "Embedded", or "Persistent" to indicate which variation of the Appender is desired.

A sample FlumeAppender configuration that is configured with a primary and a secondary agent, compresses the body, and formats the body using the RFC5424Layout:

```
1. <?xml version="1.0" encoding="UTF-8"?>
2. <Configuration status="warn" name="MyApp" packages="">
3.   <Appenders>
4.     <Flume name="eventLogger" compress="true">
5.       <Agent host="192.168.10.101" port="8800"/>
6.       <Agent host="192.168.10.102" port="8800"/>
7.       <RFC5424Layout enterpriseNumber="18060" includeMDC="true" appName="MyApp"/>
8.     </Flume>
9.   </Appenders>
10.  <Loggers>
11.    <Root level="error">
12.      <AppenderRef ref="eventLogger"/>
13.    </Root>
14.  </Loggers>
15. </Configuration>
```

A sample FlumeAppender configuration that is configured with a primary and a secondary agent, compresses the body, formats the body using the RFC5424Layout, and persists encrypted events to disk:

```
1. <?xml version="1.0" encoding="UTF-8"?>
2. <Configuration status="warn" name="MyApp" packages="">
3.   <Appenders>
4.     <Flume name="eventLogger" compress="true" type="persistent" dataDir="./logData">
5.       <Agent host="192.168.10.101" port="8800"/>
6.       <Agent host="192.168.10.102" port="8800"/>
7.       <RFC5424Layout enterpriseNumber="18060" includeMDC="true" appName="MyApp"/>
8.       <Property name="keyProvider">MySecretProvider</Property>
9.     </Flume>
10.  </Appenders>
11.  <Loggers>
12.    <Root level="error">
13.      <AppenderRef ref="eventLogger"/>
14.    </Root>
15.  </Loggers>
16. </Configuration>
```

A sample FlumeAppender configuration that is configured with a primary and a secondary agent, compresses the body, formats the body using RFC5424Layout and passes the events to an embedded Flume Agent.

```
1. <?xml version="1.0" encoding="UTF-8"?>
2. <Configuration status="warn" name="MyApp" packages="">
3.   <Appenders>
4.     <Flume name="eventLogger" compress="true" type="Embedded">
5.       <Agent host="192.168.10.101" port="8800"/>
6.       <Agent host="192.168.10.102" port="8800"/>
7.       <RFC5424Layout enterpriseNumber="18060" includeMDC="true" appName="MyApp"/>
8.     </Flume>
9.     <Console name="STDOUT">
10.      <PatternLayout pattern="%d [%p] %c %m%n"/>
11.    </Console>
12.  </Appenders>
13.  <Loggers>
14.    <Logger name="EventLogger" level="info">
15.      <AppenderRef ref="eventLogger"/>
16.    </Logger>
17.    <Root level="warn">
18.      <AppenderRef ref="STDOUT"/>
19.    </Root>
20.  </Loggers>
21. </Configuration>
```

A sample FlumeAppender configuration that is configured with a primary and a secondary agent using Flume configuration properties, compresses the body, formats the body using RFC5424Layout and passes the events to an embedded Flume Agent.

```
1. <?xml version="1.0" encoding="UTF-8"?>
2. <Configuration status="error" name="MyApp" packages="">
3.   <Appenders>
4.     <Flume name="eventLogger" compress="true" type="Embedded">
5.       <Property name="channels">file</Property>
6.       <Property name="channels.file.type">file</Property>
7.       <Property name="channels.file.checkpointDir">target/file-channel/checkpoint</Property>
8.       <Property name="channels.file.dataDirs">target/file-channel/data</Property>
9.       <Property name="sinks">agent1 agent2</Property>
10.      <Property name="sinks.agent1.channel">file</Property>
11.      <Property name="sinks.agent1.type">avro</Property>
12.      <Property name="sinks.agent1.hostname">192.168.10.101</Property>
13.      <Property name="sinks.agent1.port">8800</Property>
14.      <Property name="sinks.agent1.batch-size">100</Property>
15.      <Property name="sinks.agent2.channel">file</Property>
16.      <Property name="sinks.agent2.type">avro</Property>
17.      <Property name="sinks.agent2.hostname">192.168.10.102</Property>
18.      <Property name="sinks.agent2.port">8800</Property>
19.      <Property name="sinks.agent2.batch-size">100</Property>
20.      <Property name="sinkgroups">group1</Property>
21.      <Property name="sinkgroups.group1.sinks">agent1 agent2</Property>
22.      <Property name="sinkgroups.group1.processor.type">failover</Property>
23.      <Property name="sinkgroups.group1.processor.priority.agent1">10</Property>
24.      <Property name="sinkgroups.group1.processor.priority.agent2">5</Property>
25.      <RFC5424Layout enterpriseNumber="18060" includeMDC="true" appName="MyApp"/>
26.    </Flume>
27.    <Console name="STDOUT">
28.      <PatternLayout pattern="%d [%p] %c %m%n"/>
29.    </Console>
30.  </Appenders>
31.  <Loggers>
32.    <Logger name="EventLogger" level="info">
33.      <AppenderRef ref="eventLogger"/>
34.    </Logger>
35.    <Root level="warn">
36.      <AppenderRef ref="STDOUT"/>
37.    </Root>
38.  </Loggers>
39. </Configuration>
```

JDBCAppender

The JDBCAppender writes log events to a relational database table using standard JDBC. It can be configured to obtain JDBC connections using a JNDI DataSource or a custom factory method. Whichever approach you take, it ***must*** be backed by a connection pool. Otherwise, logging performance will suffer greatly. If batch statements are supported by the configured JDBC driver and a bufferSize is configured to be a positive number, then log events will be batched. Note that as of Log4j 2.8, there are two ways to configure log event to column mappings: the original ColumnConfig style that only allows strings and timestamps, and the new ColumnMapping plugin that uses Log4j's built-in type conversion to allow for more data types (this is the same plugin as in the Cassandra Appender).

To get off the ground quickly during development, an alternative to using a connection source based on JNDI is to use the non-pooling DriverManager connection source. This connection source uses a JDBC connection string, a user name, and a password. Optionally, you can also use properties.

JDBCAppender Parameters

Parameter Name	Type	Description
name	String	<i>Required.</i> The name of the Appender.

ignoreExceptions	boolean	The default is true, causing exceptions encountered while appending events to be internally logged and then ignored. When set to false exceptions will be propagated to the caller, instead. You must set this to false when wrapping this Appender in a FailoverAppender.		
filter	Filter	A Filter to determine if the event should be handled by this Appender. More than one Filter may be used by using a CompositeFilter.		
bufferSize	int	If an integer greater than 0, this causes the appender to buffer log events and flush whenever the buffer reaches this size.		
connectionSource	ConnectionSource	<i>Required.</i> The connections source from which database connections should be retrieved.		
tableName	String	<i>Required.</i> The name of the database table to insert log events into.		
columnConfigs	ColumnConfig[]	<i>Required (and/or columnMappings).</i> Information about the columns that log event data should be inserted into and how to insert that data. This is represented with multiple <Column> elements.		
columnMappings	ColumnMapping[]	<i>Required (and/or columnConfigs).</i> A list of column mapping configurations. Each column must specify a column name. Each column can have a conversion type specified by its fully qualified class name. By default, the conversion type is String. If the configured type is assignment-compatible with ReadOnlyStringMap (../log4j-api/apidocs/org/apache/logging/log4j/util/ReadOnlyStringMap.html) / ThreadContextMap (../log4j-api/apidocs/org/apache/logging/log4j/spi/ThreadContextMap.html) or ThreadContextStack (../log4j-api/apidocs/org/apache/logging/log4j/spi/ThreadContextStack.html), then that column will be populated with the MDC or NDC respectively (this is database-specific how they handle inserting a Map or List value). If the configured type is assignment-compatible with java.util.Date, then the log timestamp will be converted to that configured date type. If the configured type is assignment-compatible with java.sql.Clob or java.sql.NClob, then the formatted event will be set as a Clob or NClob respectively (similar to the traditional ColumnConfig plugin). If a literal attribute is given, then its value will be used as is in the INSERT query without any escaping. Otherwise, the layout or pattern specified will be converted into the configured type and stored in that column.		
immediateFail	boolean	false	When set to true, log events will not wait to try to reconnect and will fail immediately if the JDBC resources are not available. New in 2.11.2	
reconnectIntervalMillis	long	5000	If set to a value greater than 0, after an error, the JDBCDatabaseManager will attempt to reconnect to the database after waiting the specified number of milliseconds. If the reconnect fails then an exception will be thrown (which can be caught by the application if ignoreExceptions is set to false). New in 2.11.2.	

When configuring the JDBCAppender, you must specify a ConnectionSource implementation from which the Appender gets JDBC connections. You must use exactly one of the following nested elements:

- <DataSource>: Uses JNDI.
- <ConnectionFactory>: Points to a class-method pair to provide JDBC connections.
- <DriverManager>: A quick and dirty way to get off the ground, no connection pooling.
- <PoolingDriver>: Uses Apache Commons DBCP to provide connection pooling.

DataSource Parameters

Parameter Name	Type	Description
jndiName	String	<i>Required.</i> The full, prefixed JNDI name that the javax.sql.DataSource is bound to, such as java:/comp/env/jdbc/LoggingDatabase. The DataSource must be backed by a connection pool; otherwise, logging will be very slow.

ConnectionFactory Parameters

Parameter Name	Type	Description
class	Class	<i>Required.</i> The fully qualified name of a class containing a static factory method for obtaining JDBC connections.
method	Method	<i>Required.</i> The name of a static factory method for obtaining JDBC connections. This method must have no parameters and its return type must be either java.sql.Connection or DataSource. If the method returns Connections, it must obtain them from a connection pool (and they will be returned to the pool when Log4j is done with them); otherwise, logging will be very slow. If the method returns a DataSource, the DataSource will only be retrieved once, and it must be backed by a connection pool for the same reasons.

DriverManager Parameters		
Parameter Name	Type	Description
connectionString	String	<i>Required.</i> The driver-specific JDBC connection string.
userName	String	The database user name. You cannot specify both properties and a user name or password.
password	String	The database password. You cannot specify both properties and a user name or password.
driverClassName	String	The JDBC driver class name. Some old JDBC Driver can only be discovered by explicitly loading them by class name.
properties	Property[]	A list of properties. You cannot specify both properties and a user name or password.

PoolingDriver Parameters (Apache Commons DBCP)		
Parameter Name	Type	Description
DriverManager parameters	DriverManager parameters	This connection source inherits all parameter from the DriverManager connection source.
poolName	String	The pool name used to pool JDBC Connections. Defaults to example. You can use the JDBC connection string prefix jdbc:apache:commons:dbcp: followed by the pool name if you want to use a pooled connection elsewhere. For example: jdbc:apache:commons:dbcp:example.
PoolableConnectionFactory	PoolableConnectionFactory element	Defines a PoolableConnectionFactory.

PoolableConnectionFactory Parameters (Apache Commons DBCP)		
Parameter Name	Type	Description
autoCommitOnReturn	boolean	See Apache Commons DBCP PoolableConnectionFactory. (http://commons.apache.org/proper/commons-dbcp/api-2.4.0/org/apache/commons/dbcp2/PoolableConnectionFactory.html)
cacheState	boolean	See Apache Commons DBCP PoolableConnectionFactory. (http://commons.apache.org/proper/commons-dbcp/api-2.4.0/org/apache/commons/dbcp2/PoolableConnectionFactory.html)
ConnectionInitSqls	Strings	See Apache Commons DBCP PoolableConnectionFactory. (http://commons.apache.org/proper/commons-dbcp/api-2.4.0/org/apache/commons/dbcp2/PoolableConnectionFactory.html)
defaultAutoCommit	boolean	See Apache Commons DBCP PoolableConnectionFactory. (http://commons.apache.org/proper/commons-dbcp/api-2.4.0/org/apache/commons/dbcp2/PoolableConnectionFactory.html)
defaultCatalog	String	See Apache Commons DBCP PoolableConnectionFactory. (http://commons.apache.org/proper/commons-dbcp/api-2.4.0/org/apache/commons/dbcp2/PoolableConnectionFactory.html)
defaultQueryTimeoutSeconds	integer	See Apache Commons DBCP PoolableConnectionFactory. (http://commons.apache.org/proper/commons-dbcp/api-2.4.0/org/apache/commons/dbcp2/PoolableConnectionFactory.html)
defaultReadOnly	boolean	See Apache Commons DBCP PoolableConnectionFactory. (http://commons.apache.org/proper/commons-dbcp/api-2.4.0/org/apache/commons/dbcp2/PoolableConnectionFactory.html)
defaultTransactionIsolation	integer	See Apache Commons DBCP PoolableConnectionFactory. (http://commons.apache.org/proper/commons-dbcp/api-2.4.0/org/apache/commons/dbcp2/PoolableConnectionFactory.html)
disconnectionSqlCodes	Strings	See Apache Commons DBCP PoolableConnectionFactory. (http://commons.apache.org/proper/commons-dbcp/api-2.4.0/org/apache/commons/dbcp2/PoolableConnectionFactory.html)
fastFailValidation	boolean	See Apache Commons DBCP PoolableConnectionFactory. (http://commons.apache.org/proper/commons-dbcp/api-2.4.0/org/apache/commons/dbcp2/PoolableConnectionFactory.html)
maxConnLifetimeMillis	long	See Apache Commons DBCP PoolableConnectionFactory. (http://commons.apache.org/proper/commons-dbcp/api-2.4.0/org/apache/commons/dbcp2/PoolableConnectionFactory.html)
maxOpenPreparedStatements	integer	See Apache Commons DBCP PoolableConnectionFactory. (http://commons.apache.org/proper/commons-dbcp/api-2.4.0/org/apache/commons/dbcp2/PoolableConnectionFactory.html)
poolStatements	boolean	See Apache Commons DBCP PoolableConnectionFactory. (http://commons.apache.org/proper/commons-dbcp/api-2.4.0/org/apache/commons/dbcp2/PoolableConnectionFactory.html)
rollbackOnReturn	boolean	See Apache Commons DBCP PoolableConnectionFactory. (http://commons.apache.org/proper/commons-dbcp/api-2.4.0/org/apache/commons/dbcp2/PoolableConnectionFactory.html)
validationQuery	String	See Apache Commons DBCP PoolableConnectionFactory. (http://commons.apache.org/proper/commons-dbcp/api-2.4.0/org/apache/commons/dbcp2/PoolableConnectionFactory.html)
validationQueryTimeoutSeconds	integer	See Apache Commons DBCP PoolableConnectionFactory. (http://commons.apache.org/proper/commons-dbcp/api-2.4.0/org/apache/commons/dbcp2/PoolableConnectionFactory.html)

When configuring the JDBCAppender, use the nested <Column> elements to specify which columns in the table should be written to and how to write to them. The JDBCAppender uses this information to formulate a PreparedStatement to insert records without SQL injection vulnerability.

Column Parameters		
Parameter Name	Type	Description

name	String	<i>Required.</i> The name of the database column.
pattern	String	Use this attribute to insert a value or values from the log event in this column using a PatternLayout pattern. Simply specify any legal pattern in this attribute. Either this attribute, literal, or isEventTimestamp="true" must be specified, but not more than one of these.
literal	String	Use this attribute to insert a literal value in this column. The value will be included directly in the insert SQL, without any quoting (which means that if you want this to be a string, your value should contain single quotes around it like this: literal="Literal String"). This is especially useful for databases that don't support identity columns. For example, if you are using Oracle you could specify literal="NAME_OF_YOUR_SEQUENCE.NEXTVAL" to insert a unique ID in an ID column. Either this attribute, pattern, or isEventTimestamp="true" must be specified, but not more than one of these.
parameter	String	Use this attribute to insert an expression with a parameter marker '?' in this column. The value will be included directly in the insert SQL, without any quoting (which means that if you want this to be a string, your value should contain single quotes around it like this: <ColumnMapping name="instant" parameter="TIMESTAMPADD('MILLISECOND', ?, TIMESTAMP '1970-01-01')"/> You can only specify one of literal or parameter.
isEventTimestamp	boolean	Use this attribute to insert the event timestamp in this column, which should be a SQL datetime. The value will be inserted as a java.sql.Types.TIMESTAMP. Either this attribute (equal to true), pattern, or isEventTimestamp must be specified, but not more than one of these.
isUnicode	boolean	This attribute is ignored unless pattern is specified. If true or omitted (default), the value will be inserted as unicode (setNString or setNClob). Otherwise, the value will be inserted non-unicode (setString or setClob).
isClob	boolean	This attribute is ignored unless pattern is specified. Use this attribute to indicate that the column stores Character Large Objects (CLOBs). If true, the value will be inserted as a CLOB (setClob or setNClob). If false or omitted (default), the value will be inserted as a VARCHAR or NVARCHAR (setString or setNString).

ColumnMapping Parameters

Parameter Name	Type	Description
name	String	<i>Required.</i> The name of the database column.
pattern	String	Use this attribute to insert a value or values from the log event in this column using a PatternLayout pattern. Simply specify any legal pattern in this attribute. Either this attribute, literal, or isEventTimestamp="true" must be specified, but not more than one of these.
literal	String	Use this attribute to insert a literal value in this column. The value will be included directly in the insert SQL, without any quoting (which means that if you want this to be a string, your value should contain single quotes around it like this: literal="Literal String"). This is especially useful for databases that don't support identity columns. For example, if you are using Oracle you could specify literal="NAME_OF_YOUR_SEQUENCE.NEXTVAL" to insert a unique ID in an ID column. Either this attribute, pattern, or isEventTimestamp="true" must be specified, but not more than one of these.
layout	Layout	The Layout to format the LogEvent.
type	String	Conversion type name, a fully-qualified class name.

Here are a couple sample configurations for the JDBCAppender, as well as a sample factory implementation that uses Commons Pooling and Commons DBCP to pool database connections:

```
1. <?xml version="1.0" encoding="UTF-8"?>
2. <Configuration status="error">
3.   <Appenders>
4.     <JDBC name="databaseAppender" tableName="dbo.application_log">
5.       <DataSource jndiName="java:/comp/env/jdbc/LoggingDataSource" />
6.       <Column name="eventDate" isEventTimestamp="true" />
7.       <Column name="level" pattern="%level" />
8.       <Column name="logger" pattern="%logger" />
9.       <Column name="message" pattern="%message" />
10.      <Column name="exception" pattern="%ex{full}" />
11.    </JDBC>
12.  </Appenders>
13.  <Loggers>
14.    <Root level="warn">
15.      <AppenderRef ref="databaseAppender"/>
16.    </Root>
17.  </Loggers>
18. </Configuration>
```

```

1. <?xml version="1.0" encoding="UTF-8"?>
2. <Configuration status="error">
3.   <Appenders>
4.     <JDBC name="databaseAppender" tableName="LOGGING.APPLICATION_LOG">
5.       <ConnectionFactory class="net.example.db.ConnectionFactory" method="getDatabaseConnection" />
6.       <Column name="EVENT_ID" literal="LOGGING.APPLICATION_LOG_SEQUENCE.NEXTVAL" />
7.       <Column name="EVENT_DATE" isEventTimestamp="true" />
8.       <Column name="LEVEL" pattern="%level" />
9.       <Column name="LOGGER" pattern="%logger" />
10.      <Column name="MESSAGE" pattern="%message" />
11.      <Column name="THROWABLE" pattern="%ex{full}" />
12.    </JDBC>
13.  </Appenders>
14.  <Loggers>
15.    <Root level="warn">
16.      <AppenderRef ref="databaseAppender"/>
17.    </Root>
18.  </Loggers>
19. </Configuration>

```

```

1. package net.example.db;
2.
3. import java.sql.Connection;
4. import java.sql.SQLException;
5. import java.util.Properties;
6.
7. import javax.sql.DataSource;
8.
9. import org.apache.commons.dbcp.DriverManagerConnectionFactory;
10. import org.apache.commons.dbcp.PoolableConnectionFactory;
11. import org.apache.commons.dbcp.PoolableConnectionFactory;
12. import org.apache.commons.dbcp.PoolingDataSource;
13. import org.apache.commons.pool.impl.GenericObjectPool;
14.
15. public class ConnectionFactory {
16.   private static interface Singleton {
17.     final ConnectionFactory INSTANCE = new ConnectionFactory();
18.   }
19.
20.   private final DataSource dataSource;
21.
22.   private ConnectionFactory() {
23.     Properties properties = new Properties();
24.     properties.setProperty("user", "logging");
25.     properties.setProperty("password", "abc123"); // or get properties from some configuration file
26.
27.     GenericObjectPool<PoolableConnection> pool = new GenericObjectPool<PoolableConnection>();
28.     DriverManagerConnectionFactory connectionFactory = new DriverManagerConnectionFactory(
29.       "jdbc:mysql://example.org:3306/exampleDb", properties
30.     );
31.     new PoolableConnectionFactory(
32.       connectionFactory, pool, null, "SELECT 1", 3, false, false, Connection.TRANSACTION_READ_COMMITTED
33.     );
34.
35.     this.dataSource = new PoolingDataSource(pool);
36.   }
37.
38.   public static Connection getDatabaseConnection() throws SQLException {
39.     return Singleton.INSTANCE.dataSource.getConnection();
40.   }
41. }

```

This appender is MapMessage (messages.html#MapMessage)-aware.

The following configuration uses a MessageLayout to indicate that the Appender should match the keys of a MapMessage to the names of ColumnMappings when setting the values of the Appender's SQL INSERT statement. This let you insert rows for custom values in a database table based on a Log4j MapMessage instead of values from LogEvents.

```
1. <Configuration status="debug">
2.
3.   <Appenders>
4.     <Console name="STDOUT">
5.       <PatternLayout pattern="%C{1.} %m %level MDC%X%n"/>
6.     </Console>
7.     <Jdbc name="databaseAppender" tableName="dsLogEntry" ignoreExceptions="false">
8.       <DataSource jndiName="java:/comp/env/jdbc/TestDataSourceAppender" />
9.       <ColumnMapping name="Id" />
10.      <ColumnMapping name="ColumnA" />
11.      <ColumnMapping name="ColumnB" />
12.      <MessageLayout />
13.    </Jdbc>
14.  </Appenders>
15.
16.  <Loggers>
17.    <Logger name="org.apache.logging.log4j.core.appender.db" level="debug" additivity="false">
18.      <AppenderRef ref="databaseAppender" />
19.    </Logger>
20.
21.    <Root level="fatal">
22.      <AppenderRef ref="STDOUT"/>
23.    </Root>
24.  </Loggers>
25.
26. </Configuration>
```

JMS Appender

The JMS Appender sends the formatted log event to a JMS Destination.

Note that in Log4j 2.0, this appender was split into a JMSQueueAppender and a JMSTopicAppender. Starting in Log4j 2.1, these appenders were combined into the JMS Appender which makes no distinction between queues and topics. However, configurations written for 2.0 which use the <JMSQueue/> or <JMSTopic/> elements will continue to work with the new <JMS/> configuration element.

JMS Appender Parameters

Parameter Name	Type	Default	Description
factoryBindingName	String	<i>Required</i>	The name to locate in the Context that provides the ConnectionFactory (http://download.oracle.com/javase/5/api/javax/jms/ConnectionFactory.html). This can be any subinterface of ConnectionFactory as well.
factoryName	String	<i>Required</i>	The fully qualified class name that should be used to define the Initial Context Factory as defined in INITIAL_CONTEXT_FACTORY (http://download.oracle.com/javase/7/docs/api/javax/naming/Context.html#INITIAL_CONTEXT_FACTORY). If a factoryName is specified without a providerURL a warning message will be logged as this is likely to cause problems.
filter	Filter	null	A Filter to determine if the event should be handled by this Appender. More than one Filter may be used by using a CompositeFilter.
layout	Layout	<i>Required</i>	The Layout to use to format the LogEvent. <i>New since 2.9, in previous versions SerializedLayout was default.</i>
name	String	<i>Required</i>	The name of the Appender.
password	String	null	The password to use to create the JMS connection.
providerURL	String	<i>Required</i>	The URL of the provider to use as defined by PROVIDER_URL (http://download.oracle.com/javase/7/docs/api/javax/naming/Context.html#PROVIDER_URL).
destinationBindingName	String	<i>Required</i>	The name to use to locate the Destination (http://download.oracle.com/javase/5/api/javax/jms/Destination.html). This can be a Queue or Topic, and as such, the attribute names queueBindingName and topicBindingName are aliases to maintain compatibility with the Log4j 2.0 JMS appenders.
securityPrincipalName	String	null	The name of the identity of the Principal as specified by SECURITY_PRINCIPAL (http://download.oracle.com/javase/7/docs/api/javax/naming/Context.html#SECURITY_PRINCIPAL). If a securityPrincipalName is specified without securityCredentials a warning message will be logged as this is likely to cause problems.
securityCredentials	String	null	The security credentials for the principal as specified by SECURITY_CREDENTIALS (http://download.oracle.com/javase/7/docs/api/javax/naming/Context.html#SECURITY_CREDENTIALS).
ignoreExceptions	boolean	true	When true, exceptions caught while appending events are internally logged and then ignored. When false exceptions are propagated to the caller. You must set this to false when wrapping this Appender in a FailoverAppender.
immediateFail	boolean	false	When set to true, log events will not wait to try to reconnect and will fail immediately if the JMS resources are not available. New in 2.9.
reconnectIntervalMillis	long	5000	If set to a value greater than 0, after an error, the JMSManager will attempt to reconnect to the broker after waiting the specified number of milliseconds. If the reconnect fails then an exception will be thrown (which can be caught by the application if ignoreExceptions is set to false). New in 2.9.

urlPkgPrefixes	String	null	A colon-separated list of package prefixes for the class name of the factory class that will create a URL context factory as defined by URL_PKG_PREFIXES (http://download.oracle.com/javase/7/docs/api/javax/naming/Context.html#URL_PKG_PREFIXES).
userName	String	null	The user id used to create the JMS connection.

Here is a sample JMS Appender configuration:

```
1. <?xml version="1.0" encoding="UTF-8"?>
2. <Configuration status="warn" name="MyApp">
3.   <Appenders>
4.     <JMS name="jmsQueue" destinationBindingName="MyQueue"
5.       factoryBindingName="MyQueueConnectionFactory">
6.       <JsonLayout properties="true"/>
7.     </JMS>
8.   </Appenders>
9.   <Loggers>
10.    <Root level="error">
11.      <AppenderRef ref="jmsQueue"/>
12.    </Root>
13.   </Loggers>
14. </Configuration>
```

To map your Log4j MapMessages to JMS javax.jms.MapMessages, set the layout of the appender to MessageLayout with <MessageLayout /> (Since 2.9.):

```
1. <?xml version="1.0" encoding="UTF-8"?>
2. <Configuration status="warn" name="MyApp">
3.   <Appenders>
4.     <JMS name="jmsQueue" destinationBindingName="MyQueue"
5.       factoryBindingName="MyQueueConnectionFactory">
6.       <MessageLayout />
7.     </JMS>
8.   </Appenders>
9.   <Loggers>
10.    <Root level="error">
11.      <AppenderRef ref="jmsQueue"/>
12.    </Root>
13.   </Loggers>
14. </Configuration>
```

JPAAppender

As of Log4j 2.11.0, JPA support has moved from the existing module log4j-core to the new module log4j-jpa.

The JPAAppender writes log events to a relational database table using the Java Persistence API 2.1. It requires the API and a provider implementation be on the classpath. It also requires a decorated entity configured to persist to the table desired. The entity should either extend org.apache.logging.log4j.core.appender.db.jpa.BasicLogEventEntity (if you mostly want to use the default mappings) and provide at least an @Id property, or org.apache.logging.log4j.core.appender.db.jpa.AbstractLogEventWrapperEntity (if you want to significantly customize the mappings). See the Javadoc for these two classes for more information. You can also consult the source code of these two classes as an example of how to implement the entity.

JPAAppender Parameters			
Parameter Name	Type	Description	
name	String	<i>Required.</i> The name of the Appender.	
ignoreExceptions	boolean	The default is true, causing exceptions encountered while appending events to be internally logged and then ignored. When set to false exceptions will be propagated to the caller, instead. You must set this to false when wrapping this Appender in a FailoverAppender.	
filter	Filter	A Filter to determine if the event should be handled by this Appender. More than one Filter may be used by using a CompositeFilter.	
bufferSize	int	If an integer greater than 0, this causes the appender to buffer log events and flush whenever the buffer reaches this size.	
entityClassName	String	<i>Required.</i> The fully qualified name of the concrete LogEventWrapperEntity implementation that has JPA annotations mapping it to a database table.	
persistenceUnitName	String	<i>Required.</i> The name of the JPA persistence unit that should be used for persisting log events.	

Here is a sample configuration for the JPAAppender. The first XML sample is the Log4j configuration file, the second is the persistence.xml file. EclipseLink is assumed here, but any JPA 2.1 or higher provider will do. You should *always* create a *separate* persistence unit for logging, for two reasons. First, <shared-cache-mode> *must* be set to "NONE," which is usually not desired in normal JPA usage. Also, for performance reasons the logging entity should be isolated in its own persistence unit away from all other entities and you should use a non-JTA data source. Note that your persistence unit *must* also contain <class> elements for all of the org.apache.logging.log4j.core.appender.db.jpa.converter converter classes.

```

1. <?xml version="1.0" encoding="UTF-8"?>
2. <Configuration status="error">
3.   <Appenders>
4.     <JPA name="databaseAppender" persistenceUnitName="loggingPersistenceUnit"
5.       entityClassName="com.example.logging.JpaLogEntity" />
6.   </Appenders>
7.   <Loggers>
8.     <Root level="warn">
9.       <AppenderRef ref="databaseAppender"/>
10.    </Root>
11.  </Loggers>
12. </Configuration>

```

```

1. <?xml version="1.0" encoding="UTF-8"?>
2. <persistence xmlns="http://xmlns.jcp.org/xml/ns/persistence"
3.   xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
4.   xsi:schemaLocation="http://xmlns.jcp.org/xml/ns/persistence
5.     http://xmlns.jcp.org/xml/ns/persistence/persistence_2_1.xsd"
6.   version="2.1">
7.
8.   <persistence-unit name="loggingPersistenceUnit" transaction-type="RESOURCE_LOCAL">
9.     <provider>org.eclipse.persistence.jpa.PersistenceProvider</provider>
10.    <class>org.apache.logging.log4j.core.appender.db.jpa.converter.ContextMapAttributeConverter</class>
11.    <class>org.apache.logging.log4j.core.appender.db.jpa.converter.ContextMapJsonAttributeConverter</class>
12.    <class>org.apache.logging.log4j.core.appender.db.jpa.converter.ContextStackAttributeConverter</class>
13.    <class>org.apache.logging.log4j.core.appender.db.jpa.converter.ContextStackJsonAttributeConverter</class>
14.    <class>org.apache.logging.log4j.core.appender.db.jpa.converter.MarkerAttributeConverter</class>
15.    <class>org.apache.logging.log4j.core.appender.db.jpa.converter.MessageAttributeConverter</class>
16.    <class>org.apache.logging.log4j.core.appender.db.jpa.converter.StackTraceElementAttributeConverter</class>
17.    <class>org.apache.logging.log4j.core.appender.db.jpa.converter.ThrowableAttributeConverter</class>
18.    <class>com.example.logging.JpaLogEntity</class>
19.    <non-jta-data-source>jdbc/LoggingDataSource</non-jta-data-source>
20.    <shared-cache-mode>NONE</shared-cache-mode>
21.  </persistence-unit>
22.
23. </persistence>

```

```

1. package com.example.logging;
2. ...
3. @Entity
4. @Table(name="application_log", schema="dbo")
5. public class JpaLogEntity extends BasicLogEventEntity {
6.   private static final long serialVersionUID = 1L;
7.   private long id = 0L;
8.
9.   public TestEntity() {
10.    super(null);
11.  }
12.   public TestEntity(LogEvent wrappedEvent) {
13.    super(wrappedEvent);
14.  }
15.
16.   @Id
17.   @GeneratedValue(strategy = GenerationType.IDENTITY)
18.   @Column(name = "id")
19.   public long getId() {
20.    return this.id;
21.  }
22.
23.   public void setId(long id) {
24.    this.id = id;
25.  }
26.
27.   // If you want to override the mapping of any properties mapped in BasicLogEventEntity,
28.   // just override the getters and re-specify the annotations.
29. }

```

```
1. package com.example.logging;
2. ...
3. @Entity
4. @Table(name="application_log", schema="dbo")
5. public class JpaLogEntity extends AbstractLogEventWrapperEntity {
6.     private static final long serialVersionUID = 1L;
7.     private long id = 0L;
8.
9.     public TestEntity() {
10.         super(null);
11.     }
12.     public TestEntity(LogEvent wrappedEvent) {
13.         super(wrappedEvent);
14.     }
15.
16.     @Id
17.     @GeneratedValue(strategy = GenerationType.IDENTITY)
18.     @Column(name = "logEventId")
19.     public long getId() {
20.         return this.id;
21.     }
22.
23.     public void setId(long id) {
24.         this.id = id;
25.     }
26.
27.     @Override
28.     @Enumerated(EnumType.STRING)
29.     @Column(name = "level")
30.     public Level getLevel() {
31.         return this.getWrappedEvent().getLevel();
32.     }
33.
34.     @Override
35.     @Column(name = "logger")
36.     public String getLoggerName() {
37.         return this.getWrappedEvent().getLoggerName();
38.     }
39.
40.     @Override
41.     @Column(name = "message")
42.     @Convert(converter = MyMessageConverter.class)
43.     public Message getMessage() {
44.         return this.getWrappedEvent().getMessage();
45.     }
46.     ...
47. }
```

HttpAppender

The HttpAppender sends log events over HTTP. A Layout must be provided to format the LogEvent.

Will set the Content-Type header according to the layout. Additional headers can be specified with embedded Property elements.

Will wait for response from server, and throw error if no 2xx response is received.

Implemented with HttpURLConnection (<https://docs.oracle.com/javase/7/docs/api/java/net/URLConnection.html>)

HttpAppender Parameters

Parameter Name	Type	Description
name	String	The name of the Appender.
filter	Filter	A Filter to determine if the event should be handled by this Appender. More than one Filter may be used by using a CompositeFilter.
layout	Layout	The Layout to use to format the LogEvent.
Ssl	SslConfiguration	Contains the configuration for the KeyStore and TrustStore for https. Optional, uses Java runtime defaults if not specified. See SSL
verifyHostname	boolean	Whether to verify server hostname against certificate. Only valid for https. Optional, defaults to true
url	string	The URL to use. The URL scheme must be "http" or "https".
method	string	The HTTP method to use. Optional, default is "POST".
connectTimeoutMillis	integer	The connect timeout in milliseconds. Optional, default is 0 (infinite timeout).
readTimeoutMillis	integer	The socket read timeout in milliseconds. Optional, default is 0 (infinite timeout).
headers	Property[]	Additional HTTP headers to use. The values support lookups (lookups.html).

ignoreExceptions	boolean	The default is true, causing exceptions encountered while appending events to be internally logged and then ignored. When set to false exceptions will be propagated to the caller, instead. You must set this to false when wrapping this Appender in a FailoverAppender.
------------------	---------	--

Here is a sample HttpAppender configuration snippet:

```
1. <?xml version="1.0" encoding="UTF-8"?>
2.   ...
3.   <Appenders>
4.     <Http name="Http" url="https://localhost:9200/test/log4j/">
5.       <Property name="X-Java-Runtime" value="${java:runtime}" />
6.       <JsonLayout properties="true"/>
7.       <SSL>
8.         <KeyStore location="log4j2-keystore.jks" passwordEnvironmentVariable="KEYSTORE_PASSWORD"/>
9.         <TrustStore location="truststore.jks" passwordFile="${sys:user.home}/truststore.pwd"/>
10.      </SSL>
11.    </Http>
12.  </Appenders>
```

KafkaAppender

The KafkaAppender logs events to an Apache Kafka (<https://kafka.apache.org/>) topic. Each log event is sent as a Kafka record.

KafkaAppender Parameters		
Parameter Name	Type	Description
topic	String	The Kafka topic to use. Required.
key	String	The key that will be sent to Kafka with every message. Optional value defaulting to null. Any of the Lookups (./lookups.html) can be included.
filter	Filter	A Filter to determine if the event should be handled by this Appender. More than one Filter may be used by using a CompositeFilter.
layout	Layout	The Layout to use to format the LogEvent. Required, there is no default. <i>New since 2.9, in previous versions <PatternLayout pattern="%m"/> was default.</i>
name	String	The name of the Appender. Required.
ignoreExceptions	boolean	The default is true, causing exceptions encountered while appending events to be internally logged and then ignored. When set to false exceptions will be propagated to the caller, instead. You must set this to false when wrapping this Appender in a FailoverAppender.
syncSend	boolean	The default is true, causing sends to block until the record has been acknowledged by the Kafka server. When set to false sends return immediately, allowing for lower latency and significantly higher throughput. <i>New since 2.8. Be aware that this is a new addition, and it has not been extensively tested. Any failure sending to Kafka will be reported as error to StatusLogger and the log event will be dropped (the ignoreExceptions parameter will not be effective). Log events may arrive out of order to the Kafka server.</i>
properties	Property[]	You can set properties in Kafka producer properties (http://kafka.apache.org/documentation.html#producerconfigs) . You need to set the bootstrap.servers property, there are sensible default values for the others. Do not set the value.serializer nor key.serializer properties.

Here is a sample KafkaAppender configuration snippet:

```
1. <?xml version="1.0" encoding="UTF-8"?>
2.   ...
3.   <Appenders>
4.     <Kafka name="Kafka" topic="log-test">
5.       <PatternLayout pattern="%date %message"/>
6.       <Property name="bootstrap.servers">localhost:9092</Property>
7.     </Kafka>
8.   </Appenders>
```

This appender is synchronous by default and will block until the record has been acknowledged by the Kafka server, timeout for this can be set with the timeout.ms property (defaults to 30 seconds). Wrap with Async appender (<http://logging.apache.org/log4j/2.x/manual/appenders.html#AsyncAppender>) and/or set syncSend to false to log asynchronously.

This appender requires the Kafka client library (<http://kafka.apache.org/>) . Note that you need to use a version of the Kafka client library matching the Kafka server used.

*Note:*Make sure to not let org.apache.kafka log to a Kafka appender on DEBUG level, since that will cause recursive logging:

```
1. <?xml version="1.0" encoding="UTF-8"?>
2.   ...
3.   <Loggers>
4.     <Root level="DEBUG">
5.       <AppenderRef ref="Kafka"/>
6.     </Root>
7.     <Logger name="org.apache.kafka" level="INFO" /> <!-- avoid recursive logging -->
8.   </Loggers>
```

MemoryMappedFileAppender

New since 2.1. Be aware that this is a new addition, and although it has been tested on several platforms, it does not have as much track record as the other file appenders.

The `MemoryMappedFileAppender` maps a part of the specified file into memory and writes log events to this memory, relying on the operating system's virtual memory manager to synchronize the changes to the storage device. The main benefit of using memory mapped files is I/O performance. Instead of making system calls to write to disk, this appender can simply change the program's local memory, which is orders of magnitude faster. Also, in most operating systems the memory region mapped actually is the kernel's page cache (http://en.wikipedia.org/wiki/Page_cache) (file cache), meaning that no copies need to be created in user space. (TODO: performance tests that compare performance of this appender to `RandomAccessFileAppender` and `FileAppender`.)

There is some overhead with mapping a file region into memory, especially very large regions (half a gigabyte or more). The default region size is 32 MB, which should strike a reasonable balance between the frequency and the duration of remap operations. (TODO: performance test remapping various sizes.)

Similar to the `FileAppender` and the `RandomAccessFileAppender`, `MemoryMappedFileAppender` uses a `MemoryMappedFileManager` to actually perform the file I/O. While `MemoryMappedFileAppender` from different Configurations cannot be shared, the `MemoryMappedFileManagers` can be if the Manager is accessible. For example, two web applications in a servlet container can have their own configuration and safely write to the same file if `Log4j` is in a `ClassLoader` that is common to both of them.

MemoryMappedFileAppender Parameters		
Parameter Name	Type	Description
append	boolean	When true - the default, records will be appended to the end of the file. When set to false, the file will be cleared before new records are written.
fileName	String	The name of the file to write to. If the file, or any of its parent directories, do not exist, they will be created.
filters	Filter	A Filter to determine if the event should be handled by this Appender. More than one Filter may be used by using a <code>CompositeFilter</code> .
immediateFlush	boolean	When set to true, each write will be followed by a call to <code>MappedByteBuffer.force()</code> (http://docs.oracle.com/javase/7/docs/api/java/nio/MappedByteBuffer.html#force()) . This will guarantee the data is written to the storage device.
		The default for this parameter is false. This means that the data is written to the storage device even if the Java process crashes, but there may be data loss if the operating system crashes.
		Note that manually forcing a sync on every log event loses most of the performance benefits of using a memory mapped file.
		Flushing after every write is only useful when using this appender with synchronous loggers. Asynchronous loggers and appenders will automatically flush at the end of a batch of events, even if <code>immediateFlush</code> is set to false. This also guarantees the data is written to disk but is more efficient.
regionLength	int	The length of the mapped region, defaults to 32 MB (32 * 1024 * 1024 bytes). This parameter must be a value between 256 and 1,073,741,824 (1 GB or 2^30); values outside this range will be adjusted to the closest valid value. <code>Log4j</code> will round the specified value up to the nearest power of two.
layout	Layout	The Layout to use to format the <code>LogEvent</code> . If no layout is supplied the default pattern layout of "%m%n" will be used.
name	String	The name of the Appender.
ignoreExceptions	boolean	The default is true, causing exceptions encountered while appending events to be internally logged and then ignored. When set to false exceptions will be propagated to the caller, instead. You must set this to false when wrapping this Appender in a <code>FailoverAppender</code> .

Here is a sample `MemoryMappedFile` configuration:

```
1. <?xml version="1.0" encoding="UTF-8"?>
2. <Configuration status="warn" name="MyApp" packages="">
3.   <Appenders>
4.     <MemoryMappedFile name="MyFile" fileName="logs/app.log">
5.       <PatternLayout>
6.         <Pattern>%d %p %c{1.} [%t] %m%n</Pattern>
7.       </PatternLayout>
8.     </MemoryMappedFile>
9.   </Appenders>
10.  <Loggers>
11.    <Root level="error">
12.      <AppenderRef ref="MyFile"/>
13.    </Root>
14.  </Loggers>
15. </Configuration>
```

NoSQLAppender

The `NoSQLAppender` writes log events to a NoSQL database using an internal lightweight provider interface. Provider implementations currently exist for MongoDB and Apache CouchDB, and writing a custom provider is quite simple.

NoSQLAppender Parameters		
Parameter Name	Type	Description
name	String	<i>Required.</i> The name of the Appender.
ignoreExceptions	boolean	The default is true, causing exceptions encountered while appending events to be internally logged and then ignored. When set to false exceptions will be propagated to the caller, instead. You must set this to false when wrapping this Appender in a <code>FailoverAppender</code> .
filter	Filter	A Filter to determine if the event should be handled by this Appender. More than one Filter may be used by using a <code>CompositeFilter</code> .

bufferSize	int	If an integer greater than 0, this causes the appender to buffer log events and flush whenever the buffer reaches this size.
NoSqlProvider	NoSQLProvider<C extends NoSQLConnection<W, T extends NoSQLObject<W>>>	<i>Required.</i> The NoSQL provider that provides connections to the chosen NoSQL database.

You specify which NoSQL provider to use by specifying the appropriate configuration element within the <NoSql> element. The types currently supported are <MongoDb> and <CouchDb>. To create your own custom provider, read the JavaDoc for the NoSQLProvider, NoSQLConnection, and NoSQLObject classes and the documentation about creating Log4j plugins. We recommend you review the source code for the MongoDB and CouchDB providers as a guide for creating your own provider.

The following example demonstrates how log events are persisted in NoSQL databases if represented in a JSON format:

```
{
  "level": "WARN",
  "loggerName": "com.example.application.MyClass",
  "message": "Something happened that you might want to know about.",
  "source": {
    "className": "com.example.application.MyClass",
    "methodName": "exampleMethod",
    "fileName": "MyClass.java",
    "lineNumber": 81
  },
  "marker": {
    "name": "SomeMarker",
    "parent" {
      "name": "SomeParentMarker"
    }
  },
  "threadName": "Thread-1",
  "millis": 1368844166761,
  "date": "2013-05-18T02:29:26.761Z",
  "thrown": {
    "type": "java.sql.SQLException",
    "message": "Could not insert record. Connection lost.",
    "stackTrace": [
      { "className": "org.example.sql.driver.PreparedStatement$1", "methodName": "responder", "fileName": "PreparedStatement.java", "lineNumber": 1049 },
      { "className": "org.example.sql.driver.PreparedStatement", "methodName": "executeUpdate", "fileName": "PreparedStatement.java", "lineNumber": 738 },
      { "className": "com.example.application.MyClass", "methodName": "exampleMethod", "fileName": "MyClass.java", "lineNumber": 81 },
      { "className": "com.example.application.MainClass", "methodName": "main", "fileName": "MainClass.java", "lineNumber": 52 }
    ],
    "cause": {
      "type": "java.io.IOException",
      "message": "Connection lost.",
      "stackTrace": [
        { "className": "java.nio.channels.SocketChannel", "methodName": "write", "fileName": null, "lineNumber": -1 },
        { "className": "org.example.sql.driver.PreparedStatement$1", "methodName": "responder", "fileName": "PreparedStatement.java", "lineNumber": 1032 },
        { "className": "org.example.sql.driver.PreparedStatement", "methodName": "executeUpdate", "fileName": "PreparedStatement.java", "lineNumber": 738 },
        { "className": "com.example.application.MyClass", "methodName": "exampleMethod", "fileName": "MyClass.java", "lineNumber": 81 },
        { "className": "com.example.application.MainClass", "methodName": "main", "fileName": "MainClass.java", "lineNumber": 52 }
      ]
    }
  },
  "contextMap": {
    "ID": "86c3a497-4e67-4eed-9d6a-2e5797324d7b",
    "username": "JohnDoe"
  },
  "contextStack": [
    "topItem",
    "anotherItem",
    "bottomItem"
  ]
}
```

NoSQLAppender for MongoDB

Starting with Log4 2.11.0, we provide two MongoDB modules:

- log4j-mongodb2 defines the configuration element MongoDB2 matching the MongoDB Driver version 2.
- log4j-mongodb3 defines the configuration element MongoDB3 matching the MongoDB Driver version 3.

We no longer provide the module log4j-mongodb.

The module log4j-mongodb2 aliases the old configuration element MongoDB to MongoDB2.

NoSQLAppender for MongoDB 2

This section details specializations of the NoSQLAppender provider for MongoDB using the MongoDB driver version 2. The NoSQLAppender Appender writes log events to a NoSQL database using an internal lightweight provider interface.

MongoDB2 Provider Parameters		
Parameter Name	Type	Description
collectionName	String	<i>Required.</i> The name of the MongoDB collection to insert the events into.
writeConcernConstant	Field	By default, the MongoDB provider inserts records with the instructions com.mongodb.WriteConcern.ACKNOWLEDGED. Use this optional attribute to specify the name of a constant other than ACKNOWLEDGED.
writeConcernConstantClass	Class	If you specify writeConcernConstant, you can use this attribute to specify a class other than com.mongodb.WriteConcern to find the constant on (to create your own custom instructions).
factoryClassName	Class	To provide a connection to the MongoDB database, you can use this attribute and factoryMethodName to specify a class and static method to get the connection from. The method must return a com.mongodb.DB or a com.mongodb.MongoClient. If the DB is not authenticated, you must also specify a username and password. If you use the factory method for providing a connection, you must not specify the databaseName, server, or port attributes.
factoryMethodName	Method	See the documentation for attribute factoryClassName.
databaseName	String	If you do not specify a factoryClassName and factoryMethodName for providing a MongoDB connection, you must specify a MongoDB database name using this attribute. You must also specify a username and password. You can optionally also specify a server (defaults to localhost), and a port (defaults to the default MongoDB port).
server	String	See the documentation for attribute databaseName.
port	int	See the documentation for attribute databaseName.
username	String	See the documentation for attributes databaseName and factoryClassName.
password	String	See the documentation for attributes databaseName and factoryClassName.
capped	boolean	Enable support for capped collections (https://docs.mongodb.com/manual/core/capped-collections/)
collectionSize	int	Specify the size in bytes of the capped collection to use if enabled. The minimum size is 4096 bytes, and larger sizes will be increased to the nearest integer multiple of 256. See the capped collection documentation linked above for more information.

This appender is MapMessage (messages.html#MapMessage)-aware.

Here are a few sample configurations for the NoSQLAppender and MongoDB2 provider:

```
1. <?xml version="1.0" encoding="UTF-8"?>
2. <Configuration status="error">
3.   <Appenders>
4.     <NoSql name="databaseAppender">
5.       <MongoDb2 databaseName="applicationDb" collectionName="applicationLog" server="mongo.example.org"
6.         username="loggingUser" password="abc123" />
7.     </NoSql>
8.   </Appenders>
9.   <Loggers>
10.    <Root level="warn">
11.      <AppenderRef ref="databaseAppender"/>
12.    </Root>
13.  </Loggers>
14. </Configuration>
```

```
1. <?xml version="1.0" encoding="UTF-8"?>
2. <Configuration status="error">
3.   <Appenders>
4.     <NoSql name="databaseAppender">
5.       <MongoDb2 collectionName="applicationLog" factoryClassName="org.example.db.ConnectionFactory"
6.         factoryMethodName="getNewMongoClient" />
7.     </NoSql>
8.   </Appenders>
9.   <Loggers>
10.    <Root level="warn">
11.      <AppenderRef ref="databaseAppender"/>
12.    </Root>
13.  </Loggers>
14. </Configuration>
```

Starting in Log4j version 2.11.0, the provider element name is MongoDb2. The name MongoDb is now a deprecated alias for MongoDb2.

NoSQLAppender for MongoDB 3

This section details specializations of the NoSQLAppender provider for MongoDB using the MongoDB driver version 3. The NoSQLAppender Appender writes log events to a NoSQL database using an internal lightweight provider interface.

MongoDB3 Provider Parameters		
Parameter Name	Type	Description
collectionName	String	<i>Required.</i> The name of the MongoDB collection to insert the events into.

writeConcernConstant	Field	By default, the MongoDB provider inserts records with the instructions com.mongodb.WriteConcern.ACKNOWLEDGED. Use this optional attribute to specify the name of a constant other than ACKNOWLEDGED.
writeConcernConstantClass	Class	If you specify writeConcernConstant, you can use this attribute to specify a class other than com.mongodb.WriteConcern to find the constant on (to create your own custom instructions).
factoryClassName	Class	To provide a connection to the MongoDB database, you can use this attribute and factoryMethodName to specify a class and static method to get the connection from. The method must return a com.mongodb.client.MongoDatabase or a com.mongodb.MongoClient. If the com.mongodb.client.MongoDatabase is not authenticated, you must also specify a username and password. If you use the factory method for providing a connection, you must not specify the databaseName, server, or port attributes.
factoryMethodName	Method	See the documentation for attribute factoryClassName.
databaseName	String	If you do not specify a factoryClassName and factoryMethodName for providing a MongoDB connection, you must specify a MongoDB database name using this attribute. You must also specify a username and password. You can optionally also specify a server (defaults to localhost), and a port (defaults to the default MongoDB port).
server	String	See the documentation for attribute databaseName.
port	int	See the documentation for attribute databaseName.
username	String	See the documentation for attributes databaseName and factoryClassName.
password	String	See the documentation for attributes databaseName and factoryClassName.
capped	boolean	Enable support for capped collections (https://docs.mongodb.com/manual/core/capped-collections/)
collectionSize	int	Specify the size in bytes of the capped collection to use if enabled. The minimum size is 4096 bytes, and larger sizes will be increased to the nearest integer multiple of 256. See the capped collection documentation linked above for more information.

This appender is MapMessage (messages.html#MapMessage)-aware.

Here are a few sample configurations for the NoSQLAppender and MongoDB3 provider:

```
1. <?xml version="1.0" encoding="UTF-8"?>
2. <Configuration status="error">
3.   <Appenders>
4.     <NoSql name="databaseAppender">
5.       <MongoDb3 databaseName="applicationDb" collectionName="applicationLog" server="mongo.example.org"
6.         username="loggingUser" password="abc123" />
7.     </NoSql>
8.   </Appenders>
9.   <Loggers>
10.    <Root level="warn">
11.      <AppenderRef ref="databaseAppender"/>
12.    </Root>
13.   </Loggers>
14. </Configuration>
```

```
1. <?xml version="1.0" encoding="UTF-8"?>
2. <Configuration status="error">
3.   <Appenders>
4.     <NoSql name="databaseAppender">
5.       <MongoDb3 collectionName="applicationLog" factoryClassName="org.example.db.ConnectionFactory"
6.         factoryMethodName="getNewMongoClient" />
7.     </NoSql>
8.   </Appenders>
9.   <Loggers>
10.    <Root level="warn">
11.      <AppenderRef ref="databaseAppender"/>
12.    </Root>
13.   </Loggers>
14. </Configuration>
```

NoSQLAppender for Apache CouchDB

This section details specializations of the NoSQLAppender provider for CouchDB. The NoSQLAppender writes log events to a NoSQL database using an internal lightweight provider interface.

CouchDB Provider Parameters

Parameter Name	Type	Description
factoryClassName	Class	To provide a connection to the CouchDB database, you can use this attribute and factoryMethodName to specify a class and static method to get the connection from. The method must return a org.lightcouch.CouchDbClient or a org.lightcouch.CouchDbProperties. If you use the factory method for providing a connection, you must not specify the databaseName, protocol, server, port, username, or password attributes.
factoryMethodName	Method	See the documentation for attribute factoryClassName.
databaseName	String	If you do not specify a factoryClassName and factoryMethodName for providing a CouchDB connection, you must specify a CouchDB database name using this attribute. You must also specify a username and password. You can optionally also specify a protocol (defaults to http), server (defaults to localhost), and a port (defaults to 80 for http and 443 for https).

protocol	String	Must either be "http" or "https." See the documentation for attribute databaseName.
server	String	See the documentation for attribute databaseName.
port	int	See the documentation for attribute databaseName.
username	String	See the documentation for attributes databaseName.
password	String	See the documentation for attributes databaseName.

Here are a few sample configurations for the NoSQLAppender and CouchDB provider:

```
1. <?xml version="1.0" encoding="UTF-8"?>
2. <Configuration status="error">
3.   <Appenders>
4.     <NoSql name="databaseAppender">
5.       <CouchDb databaseName="applicationDb" protocol="https" server="couch.example.org"
6.         username="loggingUser" password="abc123" />
7.     </NoSql>
8.   </Appenders>
9.   <Loggers>
10.    <Root level="warn">
11.      <AppenderRef ref="databaseAppender"/>
12.    </Root>
13.  </Loggers>
14. </Configuration>
```

OutputStreamAppender

The OutputStreamAppender provides the base for many of the other Appenders such as the File and Socket appenders that write the event to an Output Stream. It cannot be directly configured. Support for immediateFlush and buffering is provided by the OutputStreamAppender. The OutputStreamAppender uses an OutputStreamManager to handle the actual I/O, allowing the stream to be shared by Appenders in multiple configurations.

RandomAccessFileAppender

The RandomAccessFileAppender is similar to the standard FileAppender except it is always buffered (this cannot be switched off) and internally it uses a ByteBuffer + RandomAccessFile instead of a BufferedOutputStream. We saw a 20-200% performance improvement compared to FileAppender with "bufferedIO=true" in our measurements (./performance.html#whichAppender). Similar to the FileAppender, RandomAccessFileAppender uses a RandomAccessFileManager to actually perform the file I/O. While RandomAccessFileAppender from different Configurations cannot be shared, the RandomAccessFileManagers can be if the Manager is accessible. For example, two web applications in a servlet container can have their own configuration and safely write to the same file if Log4j is in a ClassLoader that is common to both of them.

RandomAccessFileAppender Parameters		
Parameter Name	Type	Description
append	boolean	When true - the default, records will be appended to the end of the file. When set to false, the file will be cleared before new records are written.
fileName	String	The name of the file to write to. If the file, or any of its parent directories, do not exist, they will be created.
filters	Filter	A Filter to determine if the event should be handled by this Appender. More than one Filter may be used by using a CompositeFilter.
immediateFlush	boolean	When set to true - the default, each write will be followed by a flush. This will guarantee the data is written to disk but could impact performance. Flushing after every write is only useful when using this appender with synchronous loggers. Asynchronous loggers and appenders will automatically flush at the end of a batch of events, even if immediateFlush is set to false. This also guarantees the data is written to disk but is more efficient.
bufferSize	int	The buffer size, defaults to 262,144 bytes (256 * 1024).
layout	Layout	The Layout to use to format the LogEvent. If no layout is supplied the default pattern layout of "%m%n" will be used.
name	String	The name of the Appender.
ignoreExceptions	boolean	The default is true, causing exceptions encountered while appending events to be internally logged and then ignored. When set to false exceptions will be propagated to the caller, instead. You must set this to false when wrapping this Appender in a FailoverAppender.

Here is a sample RandomAccessFile configuration:

```
1. <?xml version="1.0" encoding="UTF-8"?>
2. <Configuration status="warn" name="MyApp" packages="">
3.   <Appenders>
4.     <RandomAccessFile name="MyFile" fileName="logs/app.log">
5.       <PatternLayout>
6.         <Pattern>%d %p %c{1.} [%t] %m%n</Pattern>
7.       </PatternLayout>
8.     </RandomAccessFile>
9.   </Appenders>
10.  <Loggers>
11.    <Root level="error">
12.      <AppenderRef ref="MyFile"/>
13.    </Root>
14.  </Loggers>
15. </Configuration>
```

RewriteAppender

The RewriteAppender allows the LogEvent to be manipulated before it is processed by another Appender. This can be used to mask sensitive information such as passwords or to inject information into each event. The RewriteAppender must be configured with a RewritePolicy (RewritePolicy). The RewriteAppender should be configured after any Appenders it references to allow it to shut down properly.

RewriteAppender Parameters

Parameter Name	Type	Description
AppenderRef	String	The name of the Appenders to call after the LogEvent has been manipulated. Multiple AppenderRef elements can be configured.
filter	Filter	A Filter to determine if the event should be handled by this Appender. More than one Filter may be used by using a CompositeFilter.
name	String	The name of the Appender.
rewritePolicy	RewritePolicy	The RewritePolicy that will manipulate the LogEvent.
ignoreExceptions	boolean	The default is true, causing exceptions encountered while appending events to be internally logged and then ignored. When set to false exceptions will be propagated to the caller, instead. You must set this to false when wrapping this Appender in a FailoverAppender.

RewritePolicy

RewritePolicy is an interface that allows implementations to inspect and possibly modify LogEvents before they are passed to Appender. RewritePolicy declares a single method named rewrite that must be implemented. The method is passed the LogEvent and can return the same event or create a new one.

MapRewritePolicy

MapRewritePolicy will evaluate LogEvents that contain a MapMessage and will add or update elements of the Map.

Parameter Name	Type	Description
mode	String	"Add" or "Update"
keyValuePair	KeyValuePair[]	An array of keys and their values.

The following configuration shows a RewriteAppender configured to add a product key and its value to the MapMessage.:

```
1. <?xml version="1.0" encoding="UTF-8"?>
2. <Configuration status="warn" name="MyApp" packages="">
3.   <Appenders>
4.     <Console name="STDOUT" target="SYSTEM_OUT">
5.       <PatternLayout pattern="%m%n"/>
6.     </Console>
7.     <Rewrite name="rewrite">
8.       <AppenderRef ref="STDOUT"/>
9.       <MapRewritePolicy mode="Add">
10.        <KeyValuePair key="product" value="TestProduct"/>
11.      </MapRewritePolicy>
12.    </Rewrite>
13.   </Appenders>
14.   <Loggers>
15.     <Root level="error">
16.       <AppenderRef ref="Rewrite"/>
17.     </Root>
18.   </Loggers>
19. </Configuration>
```

PropertiesRewritePolicy

PropertiesRewritePolicy will add properties configured on the policy to the ThreadContext Map being logged. The properties will not be added to the actual ThreadContext Map. The property values may contain variables that will be evaluated when the configuration is processed as well as when the event is logged.

Parameter Name	Type	Description
properties	Property[]	One or more Property elements to define the keys and values to be added to the ThreadContext Map.

The following configuration shows a RewriteAppender configured to add a product key and its value to the MapMessage:

```
1. <?xml version="1.0" encoding="UTF-8"?>
2. <Configuration status="warn" name="MyApp" packages="">
3.   <Appenders>
4.     <Console name="STDOUT" target="SYSTEM_OUT">
5.       <PatternLayout pattern="%m%n"/>
6.     </Console>
7.     <Rewrite name="rewrite">
8.       <AppenderRef ref="STDOUT"/>
9.       <PropertiesRewritePolicy>
10.        <Property name="user">${sys:user.name}</Property>
11.        <Property name="env">${sys:environment}</Property>
12.      </PropertiesRewritePolicy>
13.    </Rewrite>
14.  </Appenders>
15.  <Loggers>
16.    <Root level="error">
17.      <AppenderRef ref="Rewrite"/>
18.    </Root>
19.  </Loggers>
20. </Configuration>
```

LoggerNameLevelRewritePolicy

You can use this policy to make loggers in third party code less chatty by changing event levels. The LoggerNameLevelRewritePolicy will rewrite log event levels for a given logger name prefix. You configure a LoggerNameLevelRewritePolicy with a logger name prefix and a pairs of levels, where a pair defines a source level and a target level.

Parameter Name	Type	Description
logger	String	A logger name used as a prefix to test each event's logger name.
LevelPair	KeyValuePair[]	An array of keys and their values, each key is a source level, each value a target level.

The following configuration shows a RewriteAppender configured to map level INFO to DEBUG and level WARN to INFO for all loggers that start with com.foo.bar.

```
1. <?xml version="1.0" encoding="UTF-8"?>
2. <Configuration status="warn" name="MyApp">
3.   <Appenders>
4.     <Console name="STDOUT" target="SYSTEM_OUT">
5.       <PatternLayout pattern="%m%n"/>
6.     </Console>
7.     <Rewrite name="rewrite">
8.       <AppenderRef ref="STDOUT"/>
9.       <LoggerNameLevelRewritePolicy logger="com.foo.bar">
10.        <KeyValuePair key="INFO" value="DEBUG"/>
11.        <KeyValuePair key="WARN" value="INFO"/>
12.      </LoggerNameLevelRewritePolicy>
13.    </Rewrite>
14.  </Appenders>
15.  <Loggers>
16.    <Root level="error">
17.      <AppenderRef ref="Rewrite"/>
18.    </Root>
19.  </Loggers>
20. </Configuration>
```

RollingFileAppender

The RollingFileAppender is an OutputStreamAppender that writes to the File named in the fileName parameter and rolls the file over according the TriggeringPolicy and the RolloverPolicy. The RollingFileAppender uses a RollingFileManager (which extends OutputStreamManager) to actually perform the file I/O and perform the rollover. While RolloverFileAppenders from different Configurations cannot be shared, the RollingFileManagers can be if the Manager is accessible. For example, two web applications in a servlet container can have their own configuration and safely write to the same file if Log4j is in a ClassLoader that is common to both of them.

A RollingFileAppender requires a TriggeringPolicy and a RolloverStrategy. The triggering policy determines if a rollover should be performed while the RolloverStrategy defines how the rollover should be done. If no RolloverStrategy is configured, RollingFileAppender will use the DefaultRolloverStrategy. Since log4j-2.5, a custom delete action can be configured in the DefaultRolloverStrategy to run at rollover. Since 2.8 if no file name is configured then DirectWriteRolloverStrategy will be used instead of DefaultRolloverStrategy. Since log4j-2.9, a custom POSIX file attribute view action can be configured in the DefaultRolloverStrategy to run at rollover, if not defined, inherited POSIX file attribute view from the RollingFileAppender will be applied.

File locking is not supported by the RollingFileAppender.

RollingFileAppender Parameters		
Parameter Name	Type	Description
append	boolean	When true - the default, records will be appended to the end of the file. When set to false, the file will be cleared before new records are written.
bufferedIO	boolean	When true - the default, records will be written to a buffer and the data will be written to disk when the buffer is full or, if immediateFlush is set, when the record is written. File locking cannot be used with bufferedIO. Performance tests have shown that using buffered I/O significantly improves performance, even if immediateFlush is enabled.

bufferSize	int	When bufferedIO is true, this is the buffer size, the default is 8192 bytes.
createOnDemand	boolean	The appender creates the file on-demand. The appender only creates the file when a log event passes all filters and is routed to this appender. Defaults to false.
filter	Filter	A Filter to determine if the event should be handled by this Appender. More than one Filter may be used by using a CompositeFilter.
fileName	String	The name of the file to write to. If the file, or any of its parent directories, do not exist, they will be created.
filePattern	String	The pattern of the file name of the archived log file. The format of the pattern is dependent on the RolloverPolicy that is used. The DefaultRolloverPolicy will accept both a date/time pattern compatible with SimpleDateFormat (http://download.oracle.com/javase/7/docs/api/java/text/SimpleDateFormat.html) and/or a %i which represents an integer counter. The pattern also supports interpolation at runtime so any of the Lookups (such as the DateLookup (./lookups.html#DateLookup)) can be included in the pattern.
immediateFlush	boolean	<p>When set to true - the default, each write will be followed by a flush. This will guarantee the data is written to disk but could impact performance.</p> <p>Flushing after every write is only useful when using this appender with synchronous loggers. Asynchronous loggers and appenders will automatically flush at the end of a batch of events, even if immediateFlush is set to false. This also guarantees the data is written to disk but is more efficient.</p>
layout	Layout	The Layout to use to format the LogEvent. If no layout is supplied the default pattern layout of "%m%n" will be used.
name	String	The name of the Appender.
policy	TriggeringPolicy	The policy to use to determine if a rollover should occur.
strategy	RolloverStrategy	The strategy to use to determine the name and location of the archive file.
ignoreExceptions	boolean	The default is true, causing exceptions encountered while appending events to be internally logged and then ignored. When set to false exceptions will be propagated to the caller, instead. You must set this to false when wrapping this Appender in a FailoverAppender.
filePermissions	String	<p>File attribute permissions in POSIX format to apply whenever the file is created.</p> <p>Underlying files system shall support POSIX (https://docs.oracle.com/javase/7/docs/api/java/nio/file/attribute/PosixFileAttributeView.html) file attribute view.</p> <p>Examples: rw----- or rw-rw-rw- etc...</p>
fileOwner	String	<p>File owner to define whenever the file is created.</p> <p>Changing file's owner may be restricted for security reason and Operation not permitted IOException thrown. Only processes with an effective user ID equal to the user ID of the file or with appropriate privileges may change the ownership of a file if _POSIX_CHOWN_RESTRICTED (http://www.gnu.org/software/libc/manual/html_node/Options-for-Files.html) is in effect for path.</p> <p>Underlying files system shall support file owner (https://docs.oracle.com/javase/7/docs/api/java/nio/file/attribute/FileOwnerAttributeView.html) attribute view.</p>
fileGroup	String	<p>File group to define whenever the file is created.</p> <p>Underlying files system shall support POSIX (https://docs.oracle.com/javase/7/docs/api/java/nio/file/attribute/PosixFileAttributeView.html) file attribute view.</p>

Triggering Policies

Composite Triggering Policy

The CompositeTriggeringPolicy combines multiple triggering policies and returns true if any of the configured policies return true. The CompositeTriggeringPolicy is configured simply by wrapping other policies in a Policies element.

For example, the following XML fragment defines policies that rollover the log when the JVM starts, when the log size reaches twenty megabytes, and when the current date no longer matches the log's start date.

```
1. <Policies>
2.   <OnStartupTriggeringPolicy />
3.   <SizeBasedTriggeringPolicy size="20 MB" />
4.   <TimeBasedTriggeringPolicy />
5. </Policies>
```

Cron Triggering Policy

The CronTriggeringPolicy triggers rollover based on a cron expression. This policy is controlled by a timer and is asynchronous to processing log events, so it is possible that log events from the previous or next time period may appear at the beginning or end of the log file. The filePattern attribute of the Appender should contain a timestamp otherwise the target file will be overwritten on each rollover.

CronTriggeringPolicy Parameters		
Parameter Name	Type	Description
schedule	String	The cron expression. The expression is the same as what is allowed in the Quartz scheduler. See CronExpression (./log4j-core/apidocs/org/apache/logging/log4j/core/util/CronExpression.html) for a full description of the expression.
evaluateOnStartup	boolean	On startup the cron expression will be evaluated against the file's last modification timestamp. If the cron expression indicates a rollover should have occurred between that time and the current time the file will be immediately rolled over.

OnStartup Triggering Policy

The OnStartupTriggeringPolicy policy causes a rollover if the log file is older than the current JVM's start time and the minimum file size is met or exceeded.

OnStartupTriggeringPolicy Parameters		
Parameter Name	Type	Description
minSize	long	The minimum size the file must have to roll over. A size of zero will cause a roll over no matter what the file size is. The default value is 1, which will prevent rolling over an empty file.

Google App Engine note:

When running in Google App Engine, the OnStartup policy causes a rollover if the log file is older than *the time when Log4J initialized*. (Google App Engine restricts access to certain classes so Log4J cannot determine JVM start time with java.lang.management.ManagementFactory.getRuntimeMXBean().getStartTime() and falls back to Log4J initialization time instead.)

SizeBased Triggering Policy

The SizeBasedTriggeringPolicy causes a rollover once the file has reached the specified size. The size can be specified in bytes, with the suffix KB, MB or GB, for example 20MB. When combined with a time based triggering policy the file pattern must contain a %i otherwise the target file will be overwritten on every rollover as the SizeBased Triggering Policy will not cause the timestamp value in the file name to change. When used without a time based triggering policy the SizeBased Triggering Policy will cause the timestamp value to change.

TimeBased Triggering Policy

The TimeBasedTriggeringPolicy causes a rollover once the date/time pattern no longer applies to the active file. This policy accepts an interval attribute which indicates how frequently the rollover should occur based on the time pattern and a modulate boolean attribute.

TimeBasedTriggeringPolicy Parameters		
Parameter Name	Type	Description
interval	integer	How often a rollover should occur based on the most specific time unit in the date pattern. For example, with a date pattern with hours as the most specific item and an increment of 4 rollovers would occur every 4 hours. The default value is 1.
modulate	boolean	Indicates whether the interval should be adjusted to cause the next rollover to occur on the interval boundary. For example, if the item is hours, the current hour is 3 am and the interval is 4 then the first rollover will occur at 4 am and then next ones will occur at 8 am, noon, 4pm, etc.
maxRandomDelay	integer	Indicates the maximum number of seconds to randomly delay a rollover. By default, this is 0 which indicates no delay. This setting is useful on servers where multiple applications are configured to rollover log files at the same time and can spread the load of doing so across time.

Rollover Strategies

Default Rollover Strategy

The default rollover strategy accepts both a date/time pattern and an integer from the filePattern attribute specified on the RollingFileAppender itself. If the date/time pattern is present it will be replaced with the current date and time values. If the pattern contains an integer it will be incremented on each rollover. If the pattern contains both a date/time and integer in the pattern the integer will be incremented until the result of the date/time pattern changes. If the file pattern ends with ".gz", ".zip", ".bz2", ".deflate", ".pack200", or ".xz" the resulting archive will be compressed using the compression scheme that matches the suffix. The formats bzip2, Deflate, Pack200 and XZ require Apache Commons Compress (<http://commons.apache.org/proper/commons-compress/>) . In addition, XZ requires XZ for Java (<http://tukaani.org/xz/java.html>) . The pattern may also contain lookup references that can be resolved at runtime such as is shown in the example below.

The default rollover strategy supports three variations for incrementing the counter. To illustrate how it works, suppose that the min attribute is set to 1, the max attribute is set to 3, the file name is "foo.log", and the file name pattern is "foo-%i.log".

Number of rollovers	Active output target	Archived log files	Description
0	foo.log	-	All logging is going to the initial file.
1	foo.log	foo-1.log	During the first rollover foo.log is renamed to foo-1.log. A new foo.log file is created and starts being written to.
2	foo.log	foo-2.log, foo-1.log	During the second rollover foo.log is renamed to foo-2.log. A new foo.log file is created and starts being written to.
3	foo.log	foo-3.log, foo-2.log, foo-1.log	During the third rollover foo.log is renamed to foo-3.log. A new foo.log file is created and starts being written to.
4	foo.log	foo-3.log, foo-2.log, foo-1.log	In the fourth and subsequent rollovers, foo-1.log is deleted, foo-2.log is renamed to foo-1.log, foo-3.log is renamed to foo-2.log and foo.log is renamed to foo-3.log. A new foo.log file is created and starts being written to.

By way of contrast, when the fileIndex attribute is set to "min" but all the other settings are the same the "fixed window" strategy will be performed.

Number of rollovers	Active output target	Archived log files	Description
0	foo.log	-	All logging is going to the initial file.
1	foo.log	foo-1.log	During the first rollover foo.log is renamed to foo-1.log. A new foo.log file is created and starts being written to.
2	foo.log	foo-1.log, foo-2.log	During the second rollover foo-1.log is renamed to foo-2.log and foo.log is renamed to foo-1.log. A new foo.log file is created and starts being written to.

3	foo.log	foo-1.log, foo-2.log, foo-3.log	During the third rollover foo-2.log is renamed to foo-3.log, foo-1.log is renamed to foo-2.log and foo.log is renamed to foo-1.log. A new foo.log file is created and starts being written to.
4	foo.log	foo-1.log, foo-2.log, foo-3.log	In the fourth and subsequent rollovers, foo-3.log is deleted, foo-2.log is renamed to foo-3.log, foo-1.log is renamed to foo-2.log and foo.log is renamed to foo-1.log. A new foo.log file is created and starts being written to.

Finally, as of release 2.8, if the `fileIndex` attribute is set to "nomax" then the min and max values will be ignored and file numbering will increment by 1 and each rollover will have an incrementally higher value with no maximum number of files.

DefaultRolloverStrategy Parameters		
Parameter Name	Type	Description
fileIndex	String	If set to "max" (the default), files with a higher index will be newer than files with a smaller index. If set to "min", file renaming and the counter will follow the Fixed Window strategy described above.
min	integer	The minimum value of the counter. The default value is 1.
max	integer	The maximum value of the counter. Once this values is reached older archives will be deleted on subsequent rollovers. The default value is 7.
compressionLevel	integer	Sets the compression level, 0-9, where 0 = none, 1 = best speed, through 9 = best compression. Only implemented for ZIP files.
tempCompressedFilePattern	String	The pattern of the file name of the archived log file during compression.

DirectWrite Rollover Strategy

The `DirectWriteRolloverStrategy` causes log events to be written directly to files represented by the file pattern. With this strategy file renames are not performed. If the size-based triggering policy causes multiple files to be written during the specified time period they will be numbered starting at one and continually incremented until a time-based rollover occurs.

Warning: If the file pattern has a suffix indicating compression should take place the current file will not be compressed when the application is shut down. Furthermore, if the time changes such that the file pattern no longer matches the current file it will not be compressed at startup either.

DirectWriteRolloverStrategy Parameters		
Parameter Name	Type	Description
maxFiles	String	The maximum number of files to allow in the time period matching the file pattern. If the number of files is exceeded the oldest file will be deleted. If specified, the value must be greater than 1. If the value is less than zero or omitted then the number of files will not be limited.
compressionLevel	integer	Sets the compression level, 0-9, where 0 = none, 1 = best speed, through 9 = best compression. Only implemented for ZIP files.
tempCompressedFilePattern	String	The pattern of the file name of the archived log file during compression.

Below is a sample configuration that uses a `RollingFileAppender` with both the time and size based triggering policies, will create up to 7 archives on the same day (1-7) that are stored in a directory based on the current year and month, and will compress each archive using gzip:

```
1. <?xml version="1.0" encoding="UTF-8"?>
2. <Configuration status="warn" name="MyApp" packages="">
3.   <Appenders>
4.     <RollingFile name="RollingFile" fileName="logs/app.log"
5.       filePattern="logs/${date:yyyy-MM}/app-%d{MM-dd-yyyy}-%i.log.gz">
6.       <PatternLayout>
7.         <Pattern>%d %p %c{1.} [%t] %m%n</Pattern>
8.       </PatternLayout>
9.       <Policies>
10.        <TimeBasedTriggeringPolicy />
11.        <SizeBasedTriggeringPolicy size="250 MB"/>
12.      </Policies>
13.    </RollingFile>
14.  </Appenders>
15.  <Loggers>
16.    <Root level="error">
17.      <AppenderRef ref="RollingFile"/>
18.    </Root>
19.  </Loggers>
20. </Configuration>
```

This second example shows a rollover strategy that will keep up to 20 files before removing them.

```

1. <?xml version="1.0" encoding="UTF-8"?>
2. <Configuration status="warn" name="MyApp" packages="">
3.   <Appenders>
4.     <RollingFile name="RollingFile" fileName="logs/app.log"
5.       filePattern="logs/${date:yyyy-MM}/app-%d{MM-dd-yyyy}-%i.log.gz">
6.       <PatternLayout>
7.         <Pattern>%d %p %c{1.} [%t] %m%n</Pattern>
8.       </PatternLayout>
9.       <Policies>
10.        <TimeBasedTriggeringPolicy />
11.        <SizeBasedTriggeringPolicy size="250 MB"/>
12.      </Policies>
13.      <DefaultRolloverStrategy max="20"/>
14.    </RollingFile>
15.  </Appenders>
16.  <Loggers>
17.    <Root level="error">
18.      <AppenderRef ref="RollingFile"/>
19.    </Root>
20.  </Loggers>
21. </Configuration>

```

Below is a sample configuration that uses a RollingFileAppender with both the time and size based triggering policies, will create up to 7 archives on the same day (1-7) that are stored in a directory based on the current year and month, and will compress each archive using gzip and will roll every 6 hours when the hour is divisible by 6:

```

1. <?xml version="1.0" encoding="UTF-8"?>
2. <Configuration status="warn" name="MyApp" packages="">
3.   <Appenders>
4.     <RollingFile name="RollingFile" fileName="logs/app.log"
5.       filePattern="logs/${date:yyyy-MM}/app-%d{yyyy-MM-dd-HH}-%i.log.gz">
6.       <PatternLayout>
7.         <Pattern>%d %p %c{1.} [%t] %m%n</Pattern>
8.       </PatternLayout>
9.       <Policies>
10.        <TimeBasedTriggeringPolicy interval="6" modulate="true"/>
11.        <SizeBasedTriggeringPolicy size="250 MB"/>
12.      </Policies>
13.    </RollingFile>
14.  </Appenders>
15.  <Loggers>
16.    <Root level="error">
17.      <AppenderRef ref="RollingFile"/>
18.    </Root>
19.  </Loggers>
20. </Configuration>

```

This sample configuration uses a RollingFileAppender with both the cron and size based triggering policies, and writes directly to an unlimited number of archive files. The cron trigger causes a rollover every hour while the file size is limited to 250MB:

```

1. <?xml version="1.0" encoding="UTF-8"?>
2. <Configuration status="warn" name="MyApp" packages="">
3.   <Appenders>
4.     <RollingFile name="RollingFile" filePattern="logs/app-%d{yyyy-MM-dd-HH}-%i.log.gz">
5.       <PatternLayout>
6.         <Pattern>%d %p %c{1.} [%t] %m%n</Pattern>
7.       </PatternLayout>
8.       <Policies>
9.        <CronTriggeringPolicy schedule="0 0 * * * ?"/>
10.        <SizeBasedTriggeringPolicy size="250 MB"/>
11.      </Policies>
12.    </RollingFile>
13.  </Appenders>
14.  <Loggers>
15.    <Root level="error">
16.      <AppenderRef ref="RollingFile"/>
17.    </Root>
18.  </Loggers>
19. </Configuration>

```

This sample configuration is the same as the previous but limits the number of files saved each hour to 10:

```
1. <?xml version="1.0" encoding="UTF-8"?>
2. <Configuration status="warn" name="MyApp" packages="">
3.   <Appenders>
4.     <RollingFile name="RollingFile" filePattern="logs/app-%d{yyyy-MM-dd-HH}-%i.log.gz">
5.       <PatternLayout>
6.         <Pattern>%d %p %c{1.} [%t] %m%n</Pattern>
7.       </PatternLayout>
8.       <Policies>
9.         <CronTriggeringPolicy schedule="0 0 * * * ?"/>
10.        <SizeBasedTriggeringPolicy size="250 MB"/>
11.      </Policies>
12.      <DirectWriteRolloverStrategy maxFiles="10"/>
13.    </RollingFile>
14.  </Appenders>
15.  <Loggers>
16.    <Root level="error">
17.      <AppenderRef ref="RollingFile"/>
18.    </Root>
19.  </Loggers>
20. </Configuration>
```

Log Archive Retention Policy: Delete on Rollover

Log4j-2.5 introduces a Delete action that gives users more control over what files are deleted at rollover time than what was possible with the DefaultRolloverStrategy max attribute. The Delete action lets users configure one or more conditions that select the files to delete relative to a base directory.

Note that it is possible to delete any file, not just rolled over log files, so use this action with care! With the testMode parameter you can test your configuration without accidentally deleting the wrong files.

Delete Parameters

Parameter Name	Type	Description
basePath	String	<i>Required.</i> Base path from where to start scanning for files to delete.
maxDepth	int	The maximum number of levels of directories to visit. A value of 0 means that only the starting file (the base path itself) is visited, unless denied by the security manager. A value of Integer.MAX_VALUE indicates that all levels should be visited. The default is 1, meaning only the files in the specified base directory.
followLinks	boolean	Whether to follow symbolic links. Default is false.
testMode	boolean	If true, files are not deleted but instead a message is printed to the status logger (configuration.html#StatusMessages) at INFO level. Use this to do a dry run to test if the configuration works as expected. Default is false.
pathSorter	PathSorter	A plugin implementing the PathSorter (../log4j-core/apidocs/org/apache/logging/log4j/core/appender/rolling/action/PathSorter.html) interface to sort the files before selecting the files to delete. The default is to sort most recently modified files first.
pathConditions	PathCondition[]	<i>Required if no ScriptCondition is specified.</i> One or more PathCondition elements. If more than one condition is specified, they all need to accept a path before it is deleted. Conditions can be nested, in which case the inner condition(s) are evaluated only if the outer condition accepts the path. If conditions are not nested they may be evaluated in any order. Conditions can also be combined with the logical operators AND, OR and NOT by using the IfAll, IfAny and IfNot composite conditions. Users can create custom conditions or use the built-in conditions: <ul style="list-style-type: none">IfFileName - accepts files whose path (relative to the base path) matches a regular expression (https://docs.oracle.com/javase/7/docs/api/java/util/regex/Pattern.html) or a glob (https://docs.oracle.com/javase/7/docs/api/java/nio/file/FileSystem.html#getPathMatcher(java.lang.String)) .IfLastModified - accepts files that are as old as or older than the specified duration (../log4j-core/apidocs/org/apache/logging/log4j/core/appender/rolling/action/Duration.html#parseCharSequence).IfAccumulatedFileCount - accepts paths after some count threshold is exceeded during the file tree walk.IfAccumulatedFileSize - accepts paths after the accumulated file size threshold is exceeded during the file tree walk.IfAll - accepts a path if all nested conditions accept it (logical AND). Nested conditions may be evaluated in any order.IfAny - accepts a path if one of the nested conditions accept it (logical OR). Nested conditions may be evaluated in any order.IfNot - accepts a path if the nested condition does not accept it (logical NOT).
scriptCondition	ScriptCondition	<i>Required if no PathConditions are specified.</i> A ScriptCondition element specifying a script. The ScriptCondition should contain a Script, ScriptRef or ScriptFile element that specifies the logic to be executed. (See also the ScriptFilter (filters.html#Script) documentation for more examples of configuring ScriptFiles and ScriptRefs.) The script is passed a number of parameters, including a list of paths found under the base path (up to maxDepth) and must return a list with the paths to delete.

IfFileName Condition Parameters

Parameter Name	Type	Description
glob	String	<i>Required if regex not specified.</i> Matches the relative path (relative to the base path) using a limited pattern language that resembles regular expressions but with a simpler syntax (https://docs.oracle.com/javase/7/docs/api/java/nio/file/FileSystem.html#getPathMatcher(java.lang.String)) .

regex	String	<i>Required if glob not specified.</i> Matches the relative path (relative to the base path) using a regular expression as defined by the Pattern (https://docs.oracle.com/javase/7/docs/api/java/util/regex/Pattern.html) class.
nestedConditions	PathCondition[]	An optional set of nested PathConditions. If any nested conditions exist they all need to accept the file before it is deleted. Nested conditions are only evaluated if the outer condition accepts a file (if the path name matches).
IfLastModified Condition Parameters		
Parameter Name	Type	Description
age	String	<i>Required.</i> Specifies a duration (../log4j-core/apidocs/org/apache/logging/log4j/core/appender/rolling/action/Duration.html#parseCharSequence). The condition accepts files that are as old or older than the specified duration.
nestedConditions	PathCondition[]	An optional set of nested PathConditions. If any nested conditions exist they all need to accept the file before it is deleted. Nested conditions are only evaluated if the outer condition accepts a file (if the file is old enough).
IfAccumulatedFileCount Condition Parameters		
Parameter Name	Type	Description
exceeds	int	<i>Required.</i> The threshold count from which files will be deleted.
nestedConditions	PathCondition[]	An optional set of nested PathConditions. If any nested conditions exist they all need to accept the file before it is deleted. Nested conditions are only evaluated if the outer condition accepts a file (if the threshold count has been exceeded).
IfAccumulatedFileSize Condition Parameters		
Parameter Name	Type	Description
exceeds	String	<i>Required.</i> The threshold accumulated file size from which files will be deleted. The size can be specified in bytes, with the suffix KB, MB or GB, for example 20MB.
nestedConditions	PathCondition[]	An optional set of nested PathConditions. If any nested conditions exist they all need to accept the file before it is deleted. Nested conditions are only evaluated if the outer condition accepts a file (if the threshold accumulated file size has been exceeded).

Below is a sample configuration that uses a RollingFileAppender with the cron triggering policy configured to trigger every day at midnight. Archives are stored in a directory based on the current year and month. All files under the base directory that match the `"/app-*.log.gz"` glob and are 60 days old or older are deleted at rollover time.

```
1. <?xml version="1.0" encoding="UTF-8"?>
2. <Configuration status="warn" name="MyApp" packages="">
3.   <Properties>
4.     <Property name="baseDir">logs</Property>
5.   </Properties>
6.   <Appenders>
7.     <RollingFile name="RollingFile" fileName="${baseDir}/app.log"
8.       filePattern="${baseDir}/${date:yyyy-MM}/app-%d{yyyy-MM-dd}.log.gz">
9.       <PatternLayout pattern="%d %p %c{1.} [%t] %m%n" />
10.      <CronTriggeringPolicy schedule="0 0 0 * * ?"/>
11.      <DefaultRolloverStrategy>
12.        <Delete basePath="${baseDir}" maxDepth="2">
13.          <IfFileName glob="*/app-*.log.gz" />
14.          <IfLastModified age="60d" />
15.        </Delete>
16.      </DefaultRolloverStrategy>
17.    </RollingFile>
18.  </Appenders>
19.  <Loggers>
20.    <Root level="error">
21.      <AppenderRef ref="RollingFile"/>
22.    </Root>
23.  </Loggers>
24. </Configuration>
```

Below is a sample configuration that uses a RollingFileAppender with both the time and size based triggering policies, will create up to 100 archives on the same day (1-100) that are stored in a directory based on the current year and month, and will compress each archive using gzip and will roll every hour. During every rollover, this configuration will delete files that match `"/app-*.log.gz"` and are 30 days old or older, but keep the most recent 100 GB or the most recent 10 files, whichever comes first.

```
1. <?xml version="1.0" encoding="UTF-8"?>
2. <Configuration status="warn" name="MyApp" packages="">
3.   <Properties>
4.     <Property name="baseDir">logs</Property>
5.   </Properties>
6.   <Appenders>
7.     <RollingFile name="RollingFile" fileName="${baseDir}/app.log"
8.       filePattern="${baseDir}/${date:yyyy-MM}/app-%d{yyyy-MM-dd-HH}-%i.log.gz">
9.       <PatternLayout pattern="%d %p %c{1.} [%t] %m%n" />
10.      <Policies>
11.        <TimeBasedTriggeringPolicy />
12.        <SizeBasedTriggeringPolicy size="250 MB"/>
13.      </Policies>
14.      <DefaultRolloverStrategy max="100">
15.        <!--
16.          Nested conditions: the inner condition is only evaluated on files
17.          for which the outer conditions are true.
18.        -->
19.        <Delete basePath="${baseDir}" maxDepth="2">
20.          <IfFileName glob="*/app-*.log.gz">
21.            <IfLastModified age="30d">
22.              <IfAny>
23.                <IfAccumulatedFileSize exceeds="100 GB" />
24.                <IfAccumulatedFileCount exceeds="10" />
25.              </IfAny>
26.            </IfLastModified>
27.          </IfFileName>
28.        </Delete>
29.      </DefaultRolloverStrategy>
30.    </RollingFile>
31.  </Appenders>
32.  <Loggers>
33.    <Root level="error">
34.      <AppenderRef ref="RollingFile"/>
35.    </Root>
36.  </Loggers>
37. </Configuration>
```

ScriptCondition Parameters

Parameter Name	Type	Description
script	Script, ScriptFile or ScriptRef	The Script element that specifies the logic to be executed. The script is passed a list of paths found under the base path and must return the paths to delete as a <code>java.util.List<PathWithAttributes></code> (../log4j-core/apidocs/org/apache/logging/log4j/core/appender/rolling/action/PathWithAttributes.html). See also the <code>ScriptFilter</code> (filters.html#Script) documentation for an example of how <code>ScriptFiles</code> and <code>ScriptRefs</code> can be configured.

Script Parameters

Parameter Name	Type	Description
basePath	<code>java.nio.file.Path</code>	The directory from where the Delete action started scanning for files to delete. Can be used to relativize the paths in the <code>pathList</code> .
pathList	<code>java.util.List<PathWithAttributes></code> (../log4j-core/apidocs/org/apache/logging/log4j/core/appender/rolling/action/PathWithAttributes.html)>	The list of paths found under the base path up to the specified max depth, sorted most recently modified files first. The script is free to modify and return this list.
statusLogger	<code>StatusLogger</code>	The <code>StatusLogger</code> that can be used to log internal events during script execution.
configuration	<code>Configuration</code>	The <code>Configuration</code> that owns this <code>ScriptCondition</code> .
substitutor	<code>StrSubstitutor</code>	The <code>StrSubstitutor</code> used to replace lookup variables.
?	<code>String</code>	Any properties declared in the configuration.

Below is a sample configuration that uses a `RollingFileAppender` with the cron triggering policy configured to trigger every day at midnight. Archives are stored in a directory based on the current year and month. The script returns a list of rolled over files under the base directory dated Friday the 13th. The Delete action will delete all files returned by the script.

```
1. <?xml version="1.0" encoding="UTF-8"?>
2. <Configuration status="trace" name="MyApp" packages="">
3.   <Properties>
4.     <Property name="baseDir">logs</Property>
5.   </Properties>
6.   <Appenders>
7.     <RollingFile name="RollingFile" fileName="${baseDir}/app.log"
8.       filePattern="${baseDir}/${date:yyyy-MM}/app-%d{yyyyMMdd}.log.gz">
9.       <PatternLayout pattern="%d %p %c{1.} [%t] %m%n" />
10.      <CronTriggeringPolicy schedule="0 0 0 * * ?"/>
11.      <DefaultRolloverStrategy>
12.        <Delete basePath="${baseDir}" maxDepth="2">
13.          <ScriptCondition>
14.            <Script name="superstitious" language="groovy"><![CDATA[
15.              import java.nio.file.*;
16.
17.              def result = [];
18.              def pattern = ~/\d*/app-(\d*)\.log\.gz/;
19.
20.              pathList.each { pathWithAttributes ->
21.                def relative = basePath.relativeTo pathWithAttributes.path
22.                statusLogger.trace 'SCRIPT: relative path=' + relative + " (base=${basePath})";
23.
24.                // remove files dated Friday the 13th
25.
26.                def matcher = pattern.matcher(relative.toString());
27.                if (matcher.find()) {
28.                  def dateString = matcher.group(1);
29.                  def calendar = Date.parse("yyyyMMdd", dateString).toCalendar();
30.                  def friday13th = calendar.get(Calendar.DAY_OF_MONTH) == 13 \
31.                    && calendar.get(Calendar.DAY_OF_WEEK) == Calendar.FRIDAY;
32.                  if (friday13th) {
33.                    result.add pathWithAttributes;
34.                    statusLogger.trace 'SCRIPT: deleting path ' + pathWithAttributes;
35.                  }
36.                }
37.              }
38.              statusLogger.trace 'SCRIPT: returning ' + result;
39.              result;
40.            ]]>
41.          </Script>
42.        </ScriptCondition>
43.      </Delete>
44.    </DefaultRolloverStrategy>
45.  </RollingFile>
46. </Appenders>
47. <Loggers>
48.   <Root level="error">
49.     <AppenderRef ref="RollingFile"/>
50.   </Root>
51. </Loggers>
52. </Configuration>
```

Log Archive File Attribute View Policy: Custom file attribute on Rollover

Log4j-2.9 introduces a PosixViewAttribute action that gives users more control over which file attribute permissions, owner and group should be applied. The PosixViewAttribute action lets users configure one or more conditions that select the eligible files relative to a base directory.

PosixViewAttribute Parameters		
Parameter Name	Type	Description
basePath	String	<i>Required.</i> Base path from where to start scanning for files to apply attributes.
maxDepth	int	The maximum number of levels of directories to visit. A value of 0 means that only the starting file (the base path itself) is visited, unless denied by the security manager. A value of Integer.MAX_VALUE indicates that all levels should be visited. The default is 1, meaning only the files in the specified base directory.
followLinks	boolean	Whether to follow symbolic links. Default is false.
pathConditions	PathCondition[]	see DeletePathCondition
filePermissions	String	File attribute permissions in POSIX format to apply when action is executed. Underlying files system shall support POSIX (https://docs.oracle.com/javase/7/docs/api/java/nio/file/attribute/PosixFileAttributeView.html) file attribute view. Examples: rw----- or rw-rw-rw- etc...

fileOwner	String	File owner to define when action is executed. Changing file's owner may be restricted for security reason and Operation not permitted IOException thrown. Only processes with an effective user ID equal to the user ID of the file or with appropriate privileges may change the ownership of a file if <code>_POSIX_CHOWN_RESTRICTED</code> (http://www.gnu.org/software/libc/manual/html_node/Options-for-Files.html) is in effect for path. Underlying files system shall support file owner (https://docs.oracle.com/javase/7/docs/api/java/nio/file/attribute/FileOwnerAttributeView.html) attribute view.
fileGroup	String	File group to define when action is executed. Underlying files system shall support POSIX (https://docs.oracle.com/javase/7/docs/api/java/nio/file/attribute/PosixFileAttributeView.html) file attribute view.

Below is a sample configuration that uses a RollingFileAppender and defines different POSIX file attribute view for current and rolled log files.

```
1. <?xml version="1.0" encoding="UTF-8"?>
2. <Configuration status="trace" name="MyApp" packages="">
3.   <Properties>
4.     <Property name="baseDir">logs</Property>
5.   </Properties>
6.   <Appenders>
7.     <RollingFile name="RollingFile" fileName="${baseDir}/app.log"
8.       filePattern="${baseDir}/${date:yyyy-MM}/app-%d{yyyyMMdd}.log.gz"
9.       filePermissions="rw-----">
10.       <PatternLayout pattern="%d %p %c{1.} [%t] %m%n" />
11.       <CronTriggeringPolicy schedule="0 0 0 * * ?"/>
12.       <DefaultRolloverStrategy stopCustomActionsOnError="true">
13.         <PosixViewAttribute basePath="${baseDir}/${date:yyyy-MM}" filePermissions="r--r--r--">
14.           <IfFileName glob="*.gz" />
15.         </PosixViewAttribute>
16.       </DefaultRolloverStrategy>
17.     </RollingFile>
18.   </Appenders>
19.
20.   <Loggers>
21.     <Root level="error">
22.       <AppenderRef ref="RollingFile"/>
23.     </Root>
24.   </Loggers>
25.
26. </Configuration>
```

RollingRandomAccessFileAppender

The RollingRandomAccessFileAppender is similar to the standard RollingFileAppender except it is always buffered (this cannot be switched off) and internally it uses a ByteBuffer + RandomAccessFile instead of a BufferedOutputStream. We saw a 20-200% performance improvement compared to RollingFileAppender with "bufferedIO=true" in our measurements (./performance.html#whichAppender). The RollingRandomAccessFileAppender writes to the File named in the fileName parameter and rolls the file over according the TriggeringPolicy and the RolloverPolicy. Similar to the RollingFileAppender, RollingRandomAccessFileAppender uses a RollingRandomAccessFileManager to actually perform the file I/O and perform the rollover. While RollingRandomAccessFileAppender from different Configurations cannot be shared, the RollingRandomAccessFileManagers can be if the Manager is accessible. For example, two web applications in a servlet container can have their own configuration and safely write to the same file if Log4j is in a ClassLoader that is common to both of them.

A RollingRandomAccessFileAppender requires a TriggeringPolicy and a RolloverStrategy. The triggering policy determines if a rollover should be performed while the RolloverStrategy defines how the rollover should be done. If no RolloverStrategy is configured, RollingRandomAccessFileAppender will use the DefaultRolloverStrategy. Since log4j-2.5, a custom delete action can be configured in the DefaultRolloverStrategy to run at rollover.

File locking is not supported by the RollingRandomAccessFileAppender.

RollingRandomAccessFileAppender Parameters

Parameter Name	Type	Description
append	boolean	When true - the default, records will be appended to the end of the file. When set to false, the file will be cleared before new records are written.
filter	Filter	A Filter to determine if the event should be handled by this Appender. More than one Filter may be used by using a CompositeFilter.
fileName	String	The name of the file to write to. If the file, or any of its parent directories, do not exist, they will be created.
filePattern	String	The pattern of the file name of the archived log file. The format of the pattern should is dependent on the RolloverStrategu that is used. The DefaultRolloverStrategy will accept both a date/time pattern compatible with SimpleDateFormat (http://download.oracle.com/javase/7/docs/api/java/text/SimpleDateFormat.html) and/or a %i which represents an integer counter. The integer counter allows specifying a padding, like %3i for space-padding the counter to 3 digits or (usually more useful) %03i for zero-padding the counter to 3 digits. The pattern also supports interpolation at runtime so any of the Lookups (such as the DateLookup (./lookups.html#DateLookup)) can be included in the pattern.

immediateFlush	boolean	<p>When set to true - the default, each write will be followed by a flush. This will guarantee the data is written to disk but could impact performance.</p> <p>Flushing after every write is only useful when using this appender with synchronous loggers. Asynchronous loggers and appenders will automatically flush at the end of a batch of events, even if immediateFlush is set to false. This also guarantees the data is written to disk but is more efficient.</p>
bufferSize	int	The buffer size, defaults to 262,144 bytes (256 * 1024).
layout	Layout	The Layout to use to format the LogEvent. If no layout is supplied the default pattern layout of "%m%n" will be used.
name	String	The name of the Appender.
policy	TriggeringPolicy	The policy to use to determine if a rollover should occur.
strategy	RolloverStrategy	The strategy to use to determine the name and location of the archive file.
ignoreExceptions	boolean	The default is true, causing exceptions encountered while appending events to be internally logged and then ignored. When set to false exceptions will be propagated to the caller, instead. You must set this to false when wrapping this Appender in a FailoverAppender.
filePermissions	String	<p>File attribute permissions in POSIX format to apply whenever the file is created.</p> <p>Underlying files system shall support POSIX (https://docs.oracle.com/javase/7/docs/api/java/nio/file/attribute/PosixFileAttributeView.html) file attribute view.</p> <p>Examples: rw----- or rw-rw-rw- etc...</p>
fileOwner	String	<p>File owner to define whenever the file is created.</p> <p>Changing file's owner may be restricted for security reason and Operation not permitted IOException thrown. Only processes with an effective user ID equal to the user ID of the file or with appropriate privileges may change the ownership of a file if _POSIX_CHOWN_RESTRICTED (http://www.gnu.org/software/libc/manual/html_node/Options-for-Files.html) is in effect for path.</p> <p>Underlying files system shall support file owner (https://docs.oracle.com/javase/7/docs/api/java/nio/file/attribute/FileOwnerAttributeView.html) attribute view.</p>
fileGroup	String	<p>File group to define whenever the file is created.</p> <p>Underlying files system shall support POSIX (https://docs.oracle.com/javase/7/docs/api/java/nio/file/attribute/PosixFileAttributeView.html) file attribute view.</p>

Triggering Policies

See RollingFileAppender Triggering Policies.

Rollover Strategies

See RollingFileAppender Rollover Strategies.

Below is a sample configuration that uses a RollingRandomAccessFileAppender with both the time and size based triggering policies, will create up to 7 archives on the same day (1-7) that are stored in a directory based on the current year and month, and will compress each archive using gzip:

```
1. <?xml version="1.0" encoding="UTF-8"?>
2. <Configuration status="warn" name="MyApp" packages="">
3.   <Appenders>
4.     <RollingRandomAccessFile name="RollingRandomAccessFile" fileName="logs/app.log"
5.       filePattern="logs/${date:yyyy-MM}/app-%d{MM-dd-yyyy}-%i.log.gz">
6.       <PatternLayout>
7.         <Pattern>%d %p %c{1.} [%t] %m%n</Pattern>
8.       </PatternLayout>
9.       <Policies>
10.        <TimeBasedTriggeringPolicy />
11.        <SizeBasedTriggeringPolicy size="250 MB"/>
12.      </Policies>
13.    </RollingRandomAccessFile>
14.  </Appenders>
15.  <Loggers>
16.    <Root level="error">
17.      <AppenderRef ref="RollingRandomAccessFile"/>
18.    </Root>
19.  </Loggers>
20. </Configuration>
```

This second example shows a rollover strategy that will keep up to 20 files before removing them.

```
1. <?xml version="1.0" encoding="UTF-8"?>
2. <Configuration status="warn" name="MyApp" packages="">
3.   <Appenders>
4.     <RollingRandomAccessFile name="RollingRandomAccessFile" fileName="logs/app.log"
5.       filePattern="logs/${date:yyyy-MM}/app-%d{MM-dd-yyyy}-%i.log.gz">
6.       <PatternLayout>
7.         <Pattern>%d %p %c{1.} [%t] %m%n</Pattern>
8.       </PatternLayout>
9.       <Policies>
10.        <TimeBasedTriggeringPolicy />
11.        <SizeBasedTriggeringPolicy size="250 MB"/>
12.      </Policies>
13.      <DefaultRolloverStrategy max="20"/>
14.    </RollingRandomAccessFile>
15.  </Appenders>
16.  <Loggers>
17.    <Root level="error">
18.      <AppenderRef ref="RollingRandomAccessFile"/>
19.    </Root>
20.  </Loggers>
21. </Configuration>
```

Below is a sample configuration that uses a RollingRandomAccessFileAppender with both the time and size based triggering policies, will create up to 7 archives on the same day (1-7) that are stored in a directory based on the current year and month, and will compress each archive using gzip and will roll every 6 hours when the hour is divisible by 6:

```
1. <?xml version="1.0" encoding="UTF-8"?>
2. <Configuration status="warn" name="MyApp" packages="">
3.   <Appenders>
4.     <RollingRandomAccessFile name="RollingRandomAccessFile" fileName="logs/app.log"
5.       filePattern="logs/${date:yyyy-MM}/app-%d{yyyy-MM-dd-HH}-%i.log.gz">
6.       <PatternLayout>
7.         <Pattern>%d %p %c{1.} [%t] %m%n</Pattern>
8.       </PatternLayout>
9.       <Policies>
10.        <TimeBasedTriggeringPolicy interval="6" modulate="true"/>
11.        <SizeBasedTriggeringPolicy size="250 MB"/>
12.      </Policies>
13.    </RollingRandomAccessFile>
14.  </Appenders>
15.  <Loggers>
16.    <Root level="error">
17.      <AppenderRef ref="RollingRandomAccessFile"/>
18.    </Root>
19.  </Loggers>
20. </Configuration>
```

RoutingAppender

The RoutingAppender evaluates LogEvents and then routes them to a subordinate Appender. The target Appender may be an appender previously configured and may be referenced by its name or the Appender can be dynamically created as needed. The RoutingAppender should be configured after any Appenders it references to allow it to shut down properly.

You can also configure a RoutingAppender with scripts: you can run a script when the appender starts and when a route is chosen for an log event.

RoutingAppender Parameters				
Parameter Name	Type	Description		
Filter	Filter	A Filter to determine if the event should be handled by this Appender. More than one Filter may be used by using a CompositeFilter.		
name	String	The name of the Appender.		
RewritePolicy	RewritePolicy	The RewritePolicy that will manipulate the LogEvent.		
Routes	Routes	Contains one or more Route declarations to identify the criteria for choosing Appenders.		
Script	Script	This Script runs when Log4j starts the RoutingAppender and returns a String Route key to determine the default Route.		
		This script is passed the following variables:		
		RoutingAppender Script Parameters		
		Parameter Name	Type	Description
		configuration	Configuration	The active Configuration.
		staticVariables	Map	A Map shared between all script invocations for this appender instance. This is the same map passed to the Routes Script.
ignoreExceptions	boolean	The default is true, causing exceptions encountered while appending events to be internally logged and then ignored. When set to false exceptions will be propagated to the caller, instead. You must set this to false when wrapping this Appender in a FailoverAppender.		

In this example, the script causes the "ServiceWindows" route to be the default route on Windows and "ServiceOther" on all other operating systems. Note that the List Appender is one of our test appenders, any appender can be used, it is only used as a shorthand.

```
1. <?xml version="1.0" encoding="UTF-8"?>
2. <Configuration status="WARN" name="RoutingTest">
3.   <Appenders>
4.     <Routing name="Routing">
5.       <Script name="RoutingInit" language="JavaScript"><![CDATA[
6.         importPackage(java.lang);
7.         System.getProperty("os.name").search("Windows") > -1 ? "ServiceWindows" : "ServiceOther";]]>
8.     </Script>
9.     <Routes>
10.      <Route key="ServiceOther">
11.        <List name="List1" />
12.      </Route>
13.      <Route key="ServiceWindows">
14.        <List name="List2" />
15.      </Route>
16.    </Routes>
17.  </Routing>
18. </Appenders>
19. <Loggers>
20.   <Root level="error">
21.     <AppenderRef ref="Routing" />
22.   </Root>
23. </Loggers>
24. </Configuration>
```

Routes

The Routes element accepts a single attribute named "pattern". The pattern is evaluated against all the registered Lookups and the result is used to select a Route. Each Route may be configured with a key. If the key matches the result of evaluating the pattern then that Route will be selected. If no key is specified on a Route then that Route is the default. Only one Route can be configured as the default.

The Routes element may contain a Script child element. If specified, the Script is run for each log event and returns the String Route key to use.

You must specify either the pattern attribute or the Script element, but not both.

Each Route must reference an Appender. If the Route contains a ref attribute then the Route will reference an Appender that was defined in the configuration. If the Route contains an Appender definition then an Appender will be created within the context of the RoutingAppender and will be reused each time a matching Appender name is referenced through a Route.

This script is passed the following variables:

RoutingAppender Routes Script Parameters		
Parameter Name	Type	Description
configuration	Configuration	The active Configuration.
staticVariables	Map	A Map shared between all script invocations for this appender instance. This is the same map passed to the Routes Script.
logEvent	LogEvent	The log event.

In this example, the script runs for each log event and picks a route based on the presence of a Marker named "AUDIT".

```

1. <?xml version="1.0" encoding="UTF-8"?>
2. <Configuration status="WARN" name="RoutingTest">
3.   <Appenders>
4.     <Console name="STDOUT" target="SYSTEM_OUT" />
5.     <Flume name="AuditLogger" compress="true">
6.       <Agent host="192.168.10.101" port="8800"/>
7.       <Agent host="192.168.10.102" port="8800"/>
8.       <RFC5424Layout enterpriseNumber="18060" includeMDC="true" appName="MyApp"/>
9.     </Flume>
10.    <Routing name="Routing">
11.      <Routes>
12.        <Script name="RoutingInit" language="JavaScript"><![CDATA[
13.          if (logEvent.getMarker() != null && logEvent.getMarker().isInstanceOf("AUDIT")) {
14.            return "AUDIT";
15.          } else if (logEvent.getContextMap().containsKey("UserId")) {
16.            return logEvent.getContextMap().get("UserId");
17.          }
18.          return "STDOUT";]]>
19.        </Script>
20.        <Route>
21.          <RollingFile
22.            name="Rolling-${mdc:UserId}"
23.            fileName="${mdc:UserId}.log"
24.            filePattern="${mdc:UserId}_%i.log.gz">
25.            <PatternLayout>
26.              <pattern>%d %p %c{1.} [%t] %m%n</pattern>
27.            </PatternLayout>
28.            <SizeBasedTriggeringPolicy size="500" />
29.          </RollingFile>
30.        </Route>
31.        <Route ref="AuditLogger" key="AUDIT"/>
32.        <Route ref="STDOUT" key="STDOUT"/>
33.      </Routes>
34.      <IdlePurgePolicy timeToLive="15" timeUnit="minutes"/>
35.    </Routing>
36.  </Appenders>
37.  <Loggers>
38.    <Root level="error">
39.      <AppenderRef ref="Routing" />
40.    </Root>
41.  </Loggers>
42. </Configuration>

```

Purge Policy

The RoutingAppender can be configured with a PurgePolicy whose purpose is to stop and remove dormant Appenders that have been dynamically created by the RoutingAppender. Log4j currently provides the IdlePurgePolicy as the only PurgePolicy available for cleaning up the Appenders. The IdlePurgePolicy accepts 2 attributes; timeToLive, which is the number of timeUnits the Appender should survive without having any events sent to it, and timeUnit, the String representation of java.util.concurrent.TimeUnit which is used with the timeToLive attribute.

Below is a sample configuration that uses a RoutingAppender to route all Audit events to a FlumeAppender and all other events will be routed to a RollingFileAppender that captures only the specific event type. Note that the AuditAppender was predefined while the RollingFileAppenders are created as needed.

```
1. <?xml version="1.0" encoding="UTF-8"?>
2. <Configuration status="warn" name="MyApp" packages="">
3.   <Appenders>
4.     <Flume name="AuditLogger" compress="true">
5.       <Agent host="192.168.10.101" port="8800"/>
6.       <Agent host="192.168.10.102" port="8800"/>
7.       <RFC5424Layout enterpriseNumber="18060" includeMDC="true" appName="MyApp"/>
8.     </Flume>
9.     <Routing name="Routing">
10.      <Routes pattern="${sd:type}">
11.        <Route>
12.          <RollingFile name="Rolling-${sd:type}" fileName="${sd:type}.log"
13.            filePattern="${sd:type}.%i.log.gz">
14.            <PatternLayout>
15.              <pattern>%d %p %c{1.} [%t] %m%n</pattern>
16.            </PatternLayout>
17.            <SizeBasedTriggeringPolicy size="500" />
18.          </RollingFile>
19.        </Route>
20.        <Route ref="AuditLogger" key="Audit"/>
21.      </Routes>
22.      <IdlePurgePolicy timeToLive="15" timeUnit="minutes"/>
23.    </Routing>
24.  </Appenders>
25.  <Loggers>
26.    <Root level="error">
27.      <AppenderRef ref="Routing"/>
28.    </Root>
29.  </Loggers>
30. </Configuration>
```

SMTPAppender

Sends an e-mail when a specific logging event occurs, typically on errors or fatal errors.

The number of logging events delivered in this e-mail depend on the value of **BufferSize** option. The SMTPAppender keeps only the last BufferSize logging events in its cyclic buffer. This keeps memory requirements at a reasonable level while still delivering useful application context. All events in the buffer are included in the email. The buffer will contain the most recent events of level TRACE to WARN preceding the event that triggered the email.

The default behavior is to trigger sending an email whenever an ERROR or higher severity event is logged and to format it as HTML. The circumstances on when the email is sent can be controlled by setting one or more filters on the Appender. As with other Appenders, the formatting can be controlled by specifying a Layout for the Appender.

SMTPAppender Parameters

Parameter Name	Type	Description
name	String	The name of the Appender.
from	String	The email address of the sender.
replyTo	String	The comma-separated list of reply-to email addresses.
to	String	The comma-separated list of recipient email addresses.
cc	String	The comma-separated list of CC email addresses.
bcc	String	The comma-separated list of BCC email addresses.
subject	String	The subject of the email message.
bufferSize	integer	The maximum number of log events to be buffered for inclusion in the message. Defaults to 512.
layout	Layout	The Layout to use to format the LogEvent. If no layout is supplied HTML layout (layouts.html#HTMLayout) will be used.
filter	Filter	A Filter to determine if the event should be handled by this Appender. More than one Filter may be used by using a CompositeFilter.
smtpDebug	boolean	When set to true enables session debugging on STDOUT. Defaults to false.
smtpHost	String	The SMTP hostname to send to. This parameter is required.
smtpPassword	String	The password required to authenticate against the SMTP server.
smtpPort	integer	The SMTP port to send to.
smtpProtocol	String	The SMTP transport protocol (such as "smtps", defaults to "smtp").
smtpUsername	String	The username required to authenticate against the SMTP server.
ignoreExceptions	boolean	The default is true, causing exceptions encountered while appending events to be internally logged and then ignored. When set to false exceptions will be propagated to the caller, instead. You must set this to false when wrapping this Appender in a FailoverAppender.
SSL	SslConfiguration	Contains the configuration for the KeyStore and TrustStore. See SSL.

```
1. <?xml version="1.0" encoding="UTF-8"?>
2. <Configuration status="warn" name="MyApp" packages="">
3.   <Appenders>
4.     <SMTP name="Mail" subject="Error Log" to="errors@logging.apache.org" from="test@logging.apache.org"
5.       smtpHost="localhost" smtpPort="25" bufferSize="50">
6.     </SMTP>
7.   </Appenders>
8.   <Loggers>
9.     <Root level="error">
10.      <AppenderRef ref="Mail"/>
11.    </Root>
12.  </Loggers>
13. </Configuration>
```

ScriptAppenderSelector

When the configuration is built, the ScriptAppenderSelector appender calls a Script to compute an appender name. Log4j then creates one of the appender named listed under AppenderSet using the name of the ScriptAppenderSelector. After configuration, Log4j ignores the ScriptAppenderSelector. Log4j only builds the one selected appender from the configuration tree, and ignores other AppenderSet child nodes.

In the following example, the script returns the name "List2". The appender name is recorded under the name of the ScriptAppenderSelector, not the name of the selected appender, in this example, "SelectIt".

```
1. <Configuration status="WARN" name="ScriptAppenderSelectorExample">
2.   <Appenders>
3.     <ScriptAppenderSelector name="SelectIt">
4.       <Script language="JavaScript"><![CDATA[
5.         importPackage(java.lang);
6.         System.getProperty("os.name").search("Windows") > -1 ? "MyCustomWindowsAppender" : "MySyslogAppender";]]>
7.     </Script>
8.     <AppenderSet>
9.       <MyCustomWindowsAppender name="MyAppender" ... />
10.      <SyslogAppender name="MySyslog" ... />
11.    </AppenderSet>
12.  </ScriptAppenderSelector>
13. </Appenders>
14. <Loggers>
15.   <Root level="error">
16.    <AppenderRef ref="SelectIt" />
17.  </Root>
18. </Loggers>
19. </Configuration>
```

SocketAppender

The SocketAppender is an OutputStreamAppender that writes its output to a remote destination specified by a host and port. The data can be sent over either TCP or UDP and can be sent in any format. You can optionally secure communication with SSL. Note that the TCP and SSL variants write to the socket as a stream and do not expect response from the target destination. Due to limitations in the TCP protocol that means that when the target server closes its connection some log events may continue to appear to succeed until a closed connection exception is raised, causing those events to be lost. If guaranteed delivery is required a protocol that requires acknowledgements must be used.

SocketAppender Parameters

Parameter Name	Type	Description
name	String	The name of the Appender.
host	String	The name or address of the system that is listening for log events. This parameter is required. If the host name resolves to multiple IP addresses the TCP and SSL variations will fail over to the next IP address when a connection is lost.
port	integer	The port on the host that is listening for log events. This parameter must be specified.
protocol	String	"TCP" (default), "SSL" or "UDP".
SSL	SslConfiguration	Contains the configuration for the KeyStore and TrustStore. See SSL.
filter	Filter	A Filter to determine if the event should be handled by this Appender. More than one Filter may be used by using a CompositeFilter.
immediateFail	boolean	When set to true, log events will not wait to try to reconnect and will fail immediately if the socket is not available.
immediateFlush	boolean	When set to true - the default, each write will be followed by a flush. This will guarantee the data is written to disk but could impact performance.
bufferedIO	boolean	When true - the default, events are written to a buffer and the data will be written to the socket when the buffer is full or, if immediateFlush is set, when the record is written.
bufferSize	int	When bufferedIO is true, this is the buffer size, the default is 8192 bytes.
layout	Layout	The Layout to use to format the LogEvent. Required, there is no default. <i>New since 2.9, in previous versions SerializedLayout was default.</i>

reconnectionDelayMillis	integer	If set to a value greater than 0, after an error the SocketManager will attempt to reconnect to the server after waiting the specified number of milliseconds. If the reconnect fails then an exception will be thrown (which can be caught by the application if ignoreExceptions is set to false).
connectTimeoutMillis	integer	The connect timeout in milliseconds. The default is 0 (infinite timeout, like Socket.connect() methods).
ignoreExceptions	boolean	The default is true, causing exceptions encountered while appending events to be internally logged and then ignored. When set to false exceptions will be propagated to the caller, instead. You must set this to false when wrapping this Appender in a FailoverAppender.

This is an unsecured TCP configuration:

```
1. <?xml version="1.0" encoding="UTF-8"?>
2. <Configuration status="warn" name="MyApp" packages="">
3.   <Appenders>
4.     <Socket name="socket" host="localhost" port="9500">
5.       <JsonLayout properties="true"/>
6.     </Socket>
7.   </Appenders>
8.   <Loggers>
9.     <Root level="error">
10.      <AppenderRef ref="socket"/>
11.    </Root>
12.  </Loggers>
13. </Configuration>
```

This is a secured SSL configuration:

```
1. <?xml version="1.0" encoding="UTF-8"?>
2. <Configuration status="warn" name="MyApp" packages="">
3.   <Appenders>
4.     <Socket name="socket" host="localhost" port="9500">
5.       <JsonLayout properties="true"/>
6.       <SSL>
7.         <KeyStore location="log4j2-keystore.jks" passwordEnvironmentVariable="KEYSTORE_PASSWORD"/>
8.         <TrustStore location="truststore.jks" passwordFile="${sys:user.home}/truststore.pwd"/>
9.       </SSL>
10.    </Socket>
11.  </Appenders>
12.  <Loggers>
13.    <Root level="error">
14.     <AppenderRef ref="socket"/>
15.    </Root>
16.  </Loggers>
17. </Configuration>
```

SSL Configuration

Several appenders can be configured to use either a plain network connection or a Secure Socket Layer (SSL) connection. This section documents the parameters available for SSL configuration.

SSL Configuration Parameters		
Parameter Name	Type	Description
protocol	String	SSL if omitted. See also Standard names (http://docs.oracle.com/javase/7/docs/technotes/guides/security/StandardNames.html#SSLContext) .
KeyStore	KeyStore	Contains your private keys and certificates, and determines which authentication credentials to send to the remote host.
TrustStore	TrustStore	Contains the CA certificates of the remote counterparty. Determines whether the remote authentication credentials (and thus the connection) should be trusted.
verifyHostName	boolean	Toggles whether or not to perform host name verification. Defaults to false.

KeyStore

The keystore is meant to contain your private keys and certificates, and determines which authentication credentials to send to the remote host.

KeyStore Configuration Parameters		
Parameter Name	Type	Description
location	String	Path to the keystore file.
password	char[]	Plain text password to access the keystore. Cannot be combined with either passwordEnvironmentVariable or passwordFile.
passwordEnvironmentVariable	String	Name of an environment variable that holds the password. Cannot be combined with either password or passwordFile.
passwordFile	String	Path to a file that holds the password. Cannot be combined with either password or passwordEnvironmentVariable.

type	String	Optional KeyStore type, e.g. JKS, PKCS12, PKCS11, BKS, Windows-MY/Windows-ROOT, KeychainStore, etc. The default is JKS. See also Standard types (http://docs.oracle.com/javase/7/docs/technotes/guides/security/StandardNames.html#KeyStore)
keyManagerFactoryAlgorithm	String	Optional KeyManagerFactory algorithm. The default is SunX509. See also Standard algorithms (http://docs.oracle.com/javase/7/docs/technotes/guides/security/StandardNames.html#KeyManagerFactory)

TrustStore

The trust store is meant to contain the CA certificates you are willing to trust when a remote party presents its certificate. Determines whether the remote authentication credentials (and thus the connection) should be trusted.

In some cases, they can be one and the same store, although it is often better practice to use distinct stores (especially when they are file-based).

TrustStore Configuration Parameters		
Parameter Name	Type	Description
location	String	Path to the keystore file.
password	char[]	Plain text password to access the keystore. Cannot be combined with either passwordEnvironmentVariable or passwordFile.
passwordEnvironmentVariable	String	Name of an environment variable that holds the password. Cannot be combined with either password or passwordFile.
passwordFile	String	Path to a file that holds the password. Cannot be combined with either password or passwordEnvironmentVariable.
type	String	Optional KeyStore type, e.g. JKS, PKCS12, PKCS11, BKS, Windows-MY/Windows-ROOT, KeychainStore, etc. The default is JKS. See also Standard types (http://docs.oracle.com/javase/7/docs/technotes/guides/security/StandardNames.html#KeyStore)
trustManagerFactoryAlgorithm	String	Optional TrustManagerFactory algorithm. The default is SunX509. See also Standard algorithms (http://docs.oracle.com/javase/7/docs/technotes/guides/security/StandardNames.html#TrustManagerFactory)

Example

```
1.  ...
2.      <SSL>
3.      <KeyStore location="log4j2-keystore.jks" passwordEnvironmentVariable="KEYSTORE_PASSWORD"/>
4.      <TrustStore location="truststore.jks" passwordFile="${sys:user.home}/truststore.pwd"/>
5.      </SSL>
6.  ...
```

SyslogAppender

The SyslogAppender is a SocketAppender that writes its output to a remote destination specified by a host and port in a format that conforms with either the BSD Syslog format or the RFC 5424 format. The data can be sent over either TCP or UDP.

SyslogAppender Parameters		
Parameter Name	Type	Description
advertise	boolean	Indicates whether the appender should be advertised.
appName	String	The value to use as the APP-NAME in the RFC 5424 syslog record.
charset	String	The character set to use when converting the syslog String to a byte array. The String must be a valid Charset (http://download.oracle.com/javase/7/docs/api/java/nio/charset/Charset.html) . If not specified, the default system Charset will be used.
connectTimeoutMillis	integer	The connect timeout in milliseconds. The default is 0 (infinite timeout, like Socket.connect() methods).
enterpriseNumber	integer	The IANA enterprise number as described in RFC 5424 (http://tools.ietf.org/html/rfc5424#section-7.2.2)
filter	Filter	A Filter to determine if the event should be handled by this Appender. More than one Filter may be used by using a CompositeFilter.
facility	String	The facility is used to try to classify the message. The facility option must be set to one of "KERN", "USER", "MAIL", "DAEMON", "AUTH", "SYSLOG", "LPR", "NEWS", "UUCP", "CRON", "AUTHPRIV", "FTP", "NTP", "AUDIT", "ALERT", "CLOCK", "LOCAL0", "LOCAL1", "LOCAL2", "LOCAL3", "LOCAL4", "LOCAL5", "LOCAL6", or "LOCAL7". These values may be specified as upper or lower case characters.
format	String	If set to "RFC5424" the data will be formatted in accordance with RFC 5424. Otherwise, it will be formatted as a BSD Syslog record. Note that although BSD Syslog records are required to be 1024 bytes or shorter the SyslogLayout does not truncate them. The RFC5424Layout also does not truncate records since the receiver must accept records of up to 2048 bytes and may accept records that are longer.
host	String	The name or address of the system that is listening for log events. This parameter is required.
id	String	The default structured data id to use when formatting according to RFC 5424. If the LogEvent contains a StructuredDataMessage the id from the Message will be used instead of this value.
ignoreExceptions	boolean	The default is true, causing exceptions encountered while appending events to be internally logged and then ignored. When set to false exceptions will be propagated to the caller, instead. You must set this to false when wrapping this Appender in a FailoverAppender.
immediateFail	boolean	When set to true, log events will not wait to try to reconnect and will fail immediately if the socket is not available.

immediateFlush	boolean	When set to true - the default, each write will be followed by a flush. This will guarantee the data is written to disk but could impact performance.
includeMDC	boolean	Indicates whether data from the ThreadContextMap will be included in the RFC 5424 Syslog record. Defaults to true.
Layout	Layout	A custom layout which overrides the format setting.
loggerFields	List of KeyValuePairs	Allows arbitrary PatternLayout patterns to be included as specified ThreadContext fields; no default specified. To use, include a >LoggerFields< nested element, containing one or more >KeyValuePair< elements. Each >KeyValuePair< must have a key attribute, which specifies the key name which will be used to identify the field within the MDC Structured Data element, and a value attribute, which specifies the PatternLayout pattern to use as the value.
mdcExcludes	String	A comma separated list of mdc keys that should be excluded from the LogEvent. This is mutually exclusive with the mdcIncludes attribute. This attribute only applies to RFC 5424 syslog records.
mdcIncludes	String	A comma separated list of mdc keys that should be included in the FlumeEvent. Any keys in the MDC not found in the list will be excluded. This option is mutually exclusive with the mdcExcludes attribute. This attribute only applies to RFC 5424 syslog records.
mdcRequired	String	A comma separated list of mdc keys that must be present in the MDC. If a key is not present a LoggingException will be thrown. This attribute only applies to RFC 5424 syslog records.
mdcPrefix	String	A string that should be prepended to each MDC key in order to distinguish it from event attributes. The default string is "mdc:". This attribute only applies to RFC 5424 syslog records.
messageId	String	The default value to be used in the MSGID field of RFC 5424 syslog records.
name	String	The name of the Appender.
newLine	boolean	If true, a newline will be appended to the end of the syslog record. The default is false.
port	integer	The port on the host that is listening for log events. This parameter must be specified.
protocol	String	"TCP" or "UDP". This parameter is required.
SSL	SslConfiguration	Contains the configuration for the KeyStore and TrustStore. See SSL.
reconnectionDelayMillis	integer	If set to a value greater than 0, after an error the SocketManager will attempt to reconnect to the server after waiting the specified number of milliseconds. If the reconnect fails then an exception will be thrown (which can be caught by the application if ignoreExceptions is set to false).

A sample syslogAppender configuration that is configured with two SyslogAppenders, one using the BSD format and one using RFC 5424.

```
1. <?xml version="1.0" encoding="UTF-8"?>
2. <Configuration status="warn" name="MyApp" packages="">
3.   <Appenders>
4.     <Syslog name="bsd" host="localhost" port="514" protocol="TCP"/>
5.     <Syslog name="RFC5424" format="RFC5424" host="localhost" port="8514"
6.       protocol="TCP" appName="MyApp" includeMDC="true"
7.       facility="LOCAL0" enterpriseNumber="18060" newLine="true"
8.       messageId="Audit" id="App"/>
9.   </Appenders>
10.  <Loggers>
11.    <Logger name="com.mycorp" level="error">
12.      <AppenderRef ref="RFC5424"/>
13.    </Logger>
14.    <Root level="error">
15.      <AppenderRef ref="bsd"/>
16.    </Root>
17.  </Loggers>
18. </Configuration>
```

For SSL this appender writes its output to a remote destination specified by a host and port over SSL in a format that conforms with either the BSD Syslog format or the RFC 5424 format.

```
1. <?xml version="1.0" encoding="UTF-8"?>
2. <Configuration status="warn" name="MyApp" packages="">
3.   <Appenders>
4.     <TLSSyslog name="bsd" host="localhost" port="6514">
5.       <SSL>
6.         <KeyStore location="log4j2-keystore.jks" passwordEnvironmentVariable="KEYSTORE_PASSWORD"/>
7.         <TrustStore location="truststore.jks" passwordFile="${sys:user.home}/truststore.pwd"/>
8.       </SSL>
9.     </TLSSyslog>
10.  </Appenders>
11.  <Loggers>
12.    <Root level="error">
13.      <AppenderRef ref="bsd"/>
14.    </Root>
15.  </Loggers>
16. </Configuration>
```

ZeroMQ/JeroMQ Appender

The ZeroMQ appender uses the JeroMQ (<https://github.com/zeromq/jeromq>) library to send log events to one or more ZeroMQ endpoints.

This is a simple JeroMQ configuration:

```
1. <?xml version="1.0" encoding="UTF-8"?>
2. <Configuration name="JeroMQAppenderTest" status="TRACE">
3.   <Appenders>
4.     <JeroMQ name="JeroMQAppender">
5.       <Property name="endpoint">tcp://*:5556</Property>
6.       <Property name="endpoint">ipc://info-topic</Property>
7.     </JeroMQ>
8.   </Appenders>
9.   <Loggers>
10.    <Root level="info">
11.      <AppenderRef ref="JeroMQAppender"/>
12.    </Root>
13.  </Loggers>
14. </Configuration>
```

The table below describes all options. Please consult the JeroMQ and ZeroMQ documentation for details.

JeroMQ Parameters		
Parameter Name	Type	Description
name	String	The name of the Appender. Required.
Layout	layout	The Layout to use to format the LogEvent. If no layout is supplied the default pattern layout of "%m%n" will be used.
Filters	Filter	The Filter(s) of the Appender.
Properties	Property[]	One or more Property elements, named endpoint.
ignoreExceptions	boolean	If true, exceptions will be logged and suppressed. If false errors will be logged and then passed to the application.
affinity	long	The ZMQ_AFFINITY option. Defaults to 0.
backlog	long	The ZMQ_BACKLOG option. Defaults to 100.
delayAttachOnConnect	boolean	The ZMQ_DELAY_ATTACH_ON_CONNECT option. Defaults to false.
identity	byte[]	The ZMQ_IDENTITY option. Defaults to none.
ipv4Only	boolean	The ZMQ_IPV4ONLY option. Defaults to true.
linger	long	The ZMQ_LINGER option. Defaults to -1.
maxMsgSize	long	The ZMQ_MAXMSGSIZE option. Defaults to -1.
rcvHwm	long	The ZMQ_RCVHWM option. Defaults to 1000.
receiveBufferSize	long	The ZMQ_RCVBUF option. Defaults to 0.
receiveTimeOut	int	The ZMQ_RCVTIMEO option. Defaults to -1.
reconnectIVL	long	The ZMQ_RECONNECT_IVL option. Defaults to 100.
reconnectIVLMax	long	The ZMQ_RECONNECT_IVL_MAX option. Defaults to 0.
sendBufferSize	long	The ZMQ_SNDBUF option. Defaults to 0.
sendTimeOut	int	The ZMQ_SNDTIMEO option. Defaults to -1.
sndHwm	long	The ZMQ_SNDHWM option. Defaults to 1000.
tcpKeepAlive	int	The ZMQ_TCP_KEEPALIVE option. Defaults to -1.
tcpKeepAliveCount	long	The ZMQ_TCP_KEEPALIVE_CNT option. Defaults to -1.
tcpKeepAliveIdle	long	The ZMQ_TCP_KEEPALIVE_IDLE option. Defaults to -1.
tcpKeepAliveInterval	long	The ZMQ_TCP_KEEPALIVE_INTVL option. Defaults to -1.
xpubVerbose	boolean	The ZMQ_XPUB_VERBOSE option. Defaults to false.